

Interview References

1. SQL & Metrics Interview (45 min)

This technical interview will be 45 minutes, conducted via HackerRank CodePair with a member of the team. The call will consist of both technical and case study portions.

At Airbnb, SQL-type languages (Hive, Presto, etc.) are the standard for data analysis. This interview will test your ability with basic and intermediate SQL concepts such as joins, aggregations, filtering, subqueries, self-joins, and window functions.

You're encouraged to use online resources during the interview—our goal is to simulate a real work environment.

The case study assesses your understanding of product measurement and your ability to structure an analytical approach to product metrics. You may be asked to dig deeper or clarify your methods.

SQL Query Reference: Basic SELECT Examples

```
SELECT *
  FROM listings
 WHERE city = 'San Francisco';

SELECT listing_id, price, room_type
  FROM listings
 WHERE city = 'New York' AND price > 200 AND is_active = 1;

SELECT listing_id, city, price
  FROM listings
 WHERE city = 'Los Angeles' OR city = 'Miami';

SELECT *
  FROM listings
 WHERE city IN ('San Francisco', 'Seattle', 'Portland');

SELECT listing_id, city, price
  FROM listings
 WHERE price BETWEEN 100 AND 300;

SELECT listing_id, city
  FROM listings
 WHERE listing_name LIKE '%Beach%';
```

```
SELECT listing_id, city
FROM listings
WHERE description IS NULL;

SELECT DISTINCT listing_id
FROM reviews
WHERE review_text IS NOT NULL;

SELECT listing_id, city, price
FROM listings
ORDER BY price ASC;

SELECT listing_id, city, price
FROM listings
ORDER BY price DESC
LIMIT 10;

SELECT *
FROM bookings
LIMIT 20;

SELECT DISTINCT city
FROM listings;

SELECT COUNT(*) AS total_listings
FROM listings;

SELECT COUNT(*) AS la_listings
FROM listings
WHERE city = 'Los Angeles';

SELECT SUM(total_price) AS total_revenue
FROM bookings
WHERE status = 'completed';

SELECT AVG(price) AS avg_price
FROM listings;
```

```
SELECT MIN(price) AS min_price
FROM listings
WHERE is_active = 1;

SELECT MAX(price) AS max_price
FROM listings;

SELECT *
FROM bookings
WHERE booking_date >= CURRENT_DATE - INTERVAL '30 days';

SELECT *
FROM bookings
WHERE YEAR(booking_date) = 2024;

SELECT listing_id, price,
CASE
    WHEN price < 100 THEN 'Budget'
    WHEN price BETWEEN 100 AND 300 THEN 'Mid-Range'
    ELSE 'Luxury'
END AS price_category
FROM listings;

SELECT listing_id,
CONCAT(city, ', ', country) AS location
FROM listings;

SELECT listing_id, UPPER(city) AS city_upper
FROM listings;

SELECT guest_id, LOWER(email) AS email_normalized
FROM guests;

SELECT listing_id, description
FROM listings
WHERE LENGTH(description) < 50;

SELECT listing_id, ROUND(price, 2) AS price_rounded
```

```
FROM listings;

SELECT listing_id, room_type
FROM listings
WHERE room_type != 'Entire home/apt';

SELECT listing_id, city, price
FROM listings
ORDER BY city ASC, price DESC;

SELECT listing_id, COALESCE(price, 0) AS price
FROM listings;
```

Business Summary

Summary of Airbnb experimentation in the marketplace context.

Marketplace Snapshot

Experimentation System and Practices

Marketing Context and Levers to Test

Metric Wiring for Readouts

Design Patterns to Reference

Quick talk track (60–90 seconds):

[Prepare a brief summary or elevator pitch for the interview.]

If you'd like, I can help convert this to a one-pager with example interview questions for each design.

This technical interview will be 45 minutes and will be conducted via HackerRank CodePair with a member of the team. The phone interview will consist of a technical portion and a case study.

At Airbnb we use a varied set of tools for doing data analysis, but the one constant in almost all pathways is the use of SQL-type languages (Hive, Presto etc.). With that in mind, this

interview will test your knowledge of basic and intermediate SQL and may include concepts such as joining, aggregating, filtering, subqueries, self-joins, windowing functions etc.

We do not mandate that you memorize anything and you are encouraged to use any and all tools available to you (including looking up syntax on the internet). We endeavor to simulate the working environment to the best of our ability.

The case study portion of the interview will test your understanding of how to measure product development, and your ability to structure an analytical approach to solve a problem involving product metrics. The questions will be at a high level, but the team might ask you to dig a bit deeper to clearly define your approach.

1.

- i. **Q10: ORDER BY - Ascending**
- ii. : Find all listings that have guest reviews.

a. `SELECT DISTINCT`

```
listing_id  
FROM reviews  
WHERE review_text IS NOT NULL;
```

b. : Find all listings that don't have a description.

2. `SELECT listing_id, city`

```
FROM listings  
WHERE description IS NULL;
```

2. : Find all listings with names containing "Beach".

ii. `SELECT listing_id, city`

```
FROM listings  
WHERE listing_name LIKE '%Beach%';
```

ii. : Find listings with prices between \$100 and \$300 per night.

b. `SELECT listing_id, city, price`

```
FROM listings  
WHERE price BETWEEN 100 AND 300;
```

b. : Find all listings in San Francisco, Seattle, or Portland.

2. `SELECT *`

```
FROM listings  
WHERE city IN ('San Francisco', 'Seattle', 'Portland');
```

2. : Find listings in either Los Angeles or Miami.

```

ii.   SELECT listing_id, city, price
FROM listings
WHERE city = 'Los Angeles' OR city = 'Miami';
ii.   : Find all active listings in New York with price greater than $200 per
night.

b.  SELECT listing_id, price, room_type
FROM listings
WHERE city = 'New York' AND price > 200 AND is_active = 1;
b.  : Find all listings located in San Francisco.

2.   SELECT *
FROM listings
WHERE city = 'San Francisco';

```

SQL Query Reference: Basic SELECT Examples

Business Summary

Summary of Airbnb experimentation with marketplace context.

Marketplace snapshot

Topic	Key facts	Why it matters for experiments
Two-sided market	Hosts supply listings; guests demand stays and experiences. >5M hosts and >8M active listings as of Q4-2024.	Supply constraints and host behavior create interference and inventory effects. Split metrics by side and market. (The Airbnb Tech Blog)
Core metrics	Nights & Experiences Booked (N&E) and Gross Booking Value (GBV). GBV is the dollar value of	Use N&E as leading indicator; GBV leads revenue. Tie

Topic	Key facts	Why it matters for experiments
	bookings incl. host earnings, fees, cleaning, taxes; N&E counts nights and seats net of cancels. (SEC)	experiment targets to these. (SEC)
Revenue model	Service fees to guests and hosts; revenue recognized at check-in. 2024 revenue ≈ \$11.1B. (SEC)	Readouts on “revenue” lag; prefer bookings/N&E for timely calls. (SEC)
Geography	Growth accelerating in APAC and LATAM; brand pushes and local payments (e.g., Pix in Brazil) increased first-time bookers.	Stratify or randomize by market; expect heterogeneous effects and adoption frictions.

Experimentation system and practices

Area	What Airbnb published	Takeaway for interview
Philosophy & pitfalls	Classic A/B pitfalls amplified by marketplace: long booking lags, cross-device identity, host responsiveness, multiple booking flows, early-significance traps; use A/A tests and bias checks. (The Airbnb Tech Blog)	Plan fixed horizons or sequential controls; monitor booking-lag cohorts; validate assignment and logging with A/A. (The Airbnb Tech Blog)
Platform (ERF)	Internal A/B platform with Airflow/Presto; ~2.5k metrics/day and ~50k experiment-metric combos; metric hierarchy: Core, Target, Certified; precomputed dimensional cuts. (Medium)	Align analyses to Core metrics; slice by subject-level (guest/host, geo, device) and event-level dimensions. (Medium)
Ranking eval	Interleaving for search ranking: team-draft pairs, booking-based attribution, ~50x speed vs A/B	Use interleaving to triage models, then confirm with A/B on business KPIs. (Airbnb Tech)

Area	What Airbnb published	Takeaway for interview
	for ranker screening; ~82% agreement with A/B. (Airbnb Tech)	
Non-standard designs	SEO measurement at market/URL level with randomized canonical URLs and diff-in-diff with cluster-robust SEs. (Medium)	When user-level randomization is impossible, switch the unit (market/URL), use DiD or synthetic controls. (Medium)

Marketing context and levers to test

Lever	Examples they cite	Experiment/causal design notes
Acquisition mix	Brand campaigns expanded in LATAM; local rails like Pix improved conversion and first-time bookers.	Geo-level holdouts or staggered rollouts; measure first-time bookers, assisted bookings, CAC/LTV.
Product-led growth	Search ranking, pricing, payments, onboarding; 535+ features/UX upgrades driving conversion. (Airbnb Newsroom)	Use ERF; prioritize conversion-proximate KPIs that predict N&E and GBV.
Monetization	Added guest travel insurance and an extra fee for cross-currency bookings. (Airbnb Newsroom)	Price/fee tests require guardrails on conversion, cancellations, NPS, and supply elasticity.
Supply growth	Co-Host Network to unlock hosting capacity; launched across 10 countries. (Airbnb Newsroom)	Supply-side experiments: split by city or host cohorts; watch acceptance rate, availability, calendar density.

Metric wiring for readouts

Stage	Primary KPI	Guardrails / cuts
Discovery → Search	Search CTR, listing views, save rate	Device, market, new vs repeat, Rooms vs entire home. (Medium)
Consideration → Request	Message sends, request rate, IB share	Host response/accept rate, time to first response. (The Airbnb Tech Blog)
Booking	N&E, GBV, ADR	Cancellation rate, coupon/fee leakage, fraud. (SEC)
Post-booking	CS contacts, ratings, repeat rate	Cross-border vs domestic, length of stay.

Design patterns to reference

Scenario	Recommended approach
Ranking/model iteration	Interleaving → confirm via A/B on Core metrics. (Airbnb Tech)
SEO landing pages	Randomize at canonical-URL level; diff-in-diff with clustered SEs; power via simulation. (Medium)
Geo marketing	Country/region switchbacks or stepped-wedge rollouts; CUPEd/DiD for variance reduction; monitor spillovers. (Medium)
Marketplace interference	Segment by side (host/guest) and market; check capacity and acceptance bottlenecks before launch decisions. (The Airbnb Tech Blog)

Quick talk track (60–90 seconds)

- Airbnb runs an ERF-based experimentation program tied to Core, Target, and Certified metrics, with dimensional cuts by geo/device/subject. ([Medium](#))
- For ranking, they interleave to get ~50x faster read, then A/B on conversion KPIs. ([Airbnb Tech](#))
- For SEO and other externally mediated surfaces, they randomize at the market/URL level and estimate effects via diff-in-diff with clustered errors. ([Medium](#))
- Top-line company metrics are N&E and GBV; revenue is fee-based and recognized at check-in, so near-term reads use bookings, not revenue. ([SEC](#))
- Current growth focus includes expansion markets, brand investment, and localized payments that raise first-time bookers.

If you want, I can convert this into a one-pager with example interview prompts tied to each design.

2. : Filtering with WHERE

3. : AND operator

4. : OR logic

5. : IN clause

6. : Range filtering

7. : Pattern matching

8. : IS NULL

9. : IS NOT NULL

10. : Sorting

Question 2	Concept	Difficulty	Business Purpose
Find all listings located in San Francisco.	Filtering with WHERE	Easy	Find listings in specific market for regional analysis

SELECT *

```
FROM listings  
WHERE city = 'San Francisco';
```

Question 3	Concept	Difficulty	Business Purpose

Find all active listings in New York with price > \$200.	AND operator	Easy	Identify premium active listings for marketing campaigns
--	--------------	------	--

```
SELECT listing_id, price, room_type
  FROM listings
 WHERE city = 'New York' AND price > 200 AND is_active = 1;
```

Question 4	Concept	Difficulty	Business Purpose
Find listings in either Los Angeles or Miami.	OR logic	Easy	Analyze multi-market presence

```
SELECT listing_id, city, price
  FROM listings
 WHERE city = 'Los Angeles' OR city = 'Miami';
```

Question 5	Concept	Difficulty	Business Purpose
Find all listings in San Francisco, Seattle, or Portland.	IN clause	Easy	Multi-city market analysis

```
SELECT *
  FROM listings
 WHERE city IN ('San Francisco', 'Seattle', 'Portland');
```

Question 6	Concept	Difficulty	Business Purpose
Find listings with prices between \$100 and \$300.	Range filtering	Easy	Identify mid-range pricing segments

```
SELECT listing_id, city, price
  FROM listings
```

```
WHERE price BETWEEN 100 AND 300;
```

Question 7	Concept	Difficulty	Business Purpose
Find all listings with names containing "Beach".	Pattern match	Easy	Search listings by name patterns for category

```
SELECT listing_id, city  
FROM listings  
WHERE listing_name LIKE '%Beach%';
```

Question 8	Concept	Difficulty	Business Purpose
Find all listings that don't have a description.	IS NULL	Easy	Data quality check for incomplete profiles

```
SELECT listing_id, city  
FROM listings  
WHERE description IS NULL;
```

Question 9	Concept	Difficulty	Business Purpose
Find all listings that have guest reviews.	IS NOT NULL	Easy	Identify complete listings for quality scoring

```
SELECT DISTINCT listing_id  
FROM reviews  
WHERE review_text IS NOT NULL;
```

Question 10	Concept	Difficulty	Business Purpose
List all listings ordered by price from	Sorting	Easy	Price analysis from low to high

low to high.			
--------------	--	--	--

```
SELECT listing_id, city, price
    FROM listings
   ORDER BY price ASC;
```

Question 11	Concept	Difficulty	Business Purpose
List the top 10 most expensive listings.	Sorting DESC	Easy	Identify premium inventory

```
SELECT listing_id, city, price
    FROM listings
   ORDER BY price DESC
   LIMIT 10;
```

Question 12	Concept	Difficulty	Business Purpose
Retrieve the first 20 bookings from the bookings table.	Result limiting	Easy	Sample data for quick analysis

```
SELECT *
    FROM bookings
   LIMIT 20;
```

Question 13	Concept	Difficulty	Business Purpose
Find all unique cities where Airbnb has listings.	Unique values	Easy	Identify all markets we operate in

```
SELECT DISTINCT city
    FROM listings;
```

Question 14	Concept	Difficulty	Business Purpose
-------------	---------	------------	------------------

Count the total number of listings.	COUNT function	Easy	Total inventory metrics

```
SELECT COUNT(*) AS total_listings
    FROM listings;
```

Question 15	Concept	Difficulty	Business Purpose
Count how many listings are in Los Angeles.	COUNT with WHERE	Easy	Market-specific inventory

```
SELECT COUNT(*) AS la_listings
    FROM listings
    WHERE city = 'Los Angeles' ;
```

Question 16	Concept	Difficulty	Business Purpose
Calculate the total revenue from all completed bookings.	Aggregation - SUM	Easy	Calculate total revenue

```
SELECT SUM(total_price) AS total_revenue
    FROM bookings
    WHERE status = 'completed' ;
```

Question 17	Concept	Difficulty	Business Purpose
Find the average nightly price of all listings.	Aggregation - AVG	Easy	Average pricing analysis

```
SELECT AVG(price) AS avg_price
    FROM listings;
```

Question 18	Concept	Difficulty	Business Purpose
Find the minimum price among all listings.	Aggregation - MIN	Easy	Identify entry-level pricing

```
SELECT MIN(price) AS min_price
  FROM listings
 WHERE is_active = 1;
```

Question 19	Concept	Difficulty	Business Purpose
Find the maximum price among all listings.	Aggregation - MAX	Easy	Identify luxury segment pricing

```
SELECT MAX(price) AS max_price
  FROM listings;
```

Question 20	Concept	Difficulty	Business Purpose
Find all bookings made in the last 30 days.	Date comparison	Easy	Recent bookings analysis

```
SELECT *
  FROM bookings
 WHERE booking_date >= CURRENT_DATE - INTERVAL '30 days';
```

Question 21	Concept	Difficulty	Business Purpose
Find all bookings made in 2024.	Date extraction	Easy	Year-over-year analysis

```
SELECT *
```

```

FROM bookings
WHERE YEAR(booking_date) = 2024;

```

Question 22	Concept	Difficulty	Business Purpose
Categorize listings as 'Budget', 'Mid-Range', or 'Luxury' by price.	Conditional logic	Easy	Categorize listings by price tier

```

SELECT listing_id, price,
CASE
    WHEN price < 100 THEN 'Budget'
    WHEN price BETWEEN 100 AND 300 THEN 'Mid-Range'
    ELSE 'Luxury'
END AS price_category
FROM listings;

```

Question 23	Concept	Difficulty	Business Purpose
Concatenate city and country for each listing.	String concatenation	Easy	Create display names for reporting

```

SELECT listing_id,
CONCAT(city, ', ', country) AS location
FROM listings;

```

Question 24	Concept	Difficulty	Business Purpose
Convert all city names to uppercase.	String manipulation	Easy	Standardize city names for grouping

```

SELECT listing_id, UPPER(city) AS city_upper
FROM listings;

```

Question 25	Concept	Difficulty	Business Purpose

Convert all guest emails to lowercase.	String manipulation	Easy	Normalize email addresses

```
SELECT guest_id, LOWER(email) AS email_normalized
FROM guests;
```

Question 26	Concept	Difficulty	Business Purpose
Find listings with description length less than 50 characters.	String length	Easy	Identify incomplete descriptions

```
SELECT listing_id, description
FROM listings
WHERE LENGTH(description) < 50;
```

Question 27	Concept	Difficulty	Business Purpose
Round all listing prices to 2 decimal places.	Number formatting	Easy	Display prices with 2 decimals

```
SELECT listing_id, ROUND(price, 2) AS price_rounded
FROM listings;
```

Question 28	Concept	Difficulty	Business Purpose
Find all listings that are NOT entire home/apt.	Negation	Easy	Exclude specific segments

```
SELECT listing_id, room_type
FROM listings
```

```
WHERE room_type != 'Entire home/apt';
```

Question 29	Concept	Difficulty	Business Purpose
List listings ordered by city, then by price within each city.	Multi-column sort	Easy	Organize listings by location/price

```
SELECT listing_id, city, price  
FROM listings  
ORDER BY city ASC, price DESC;
```

Question 30	Concept	Difficulty	Business Purpose
Replace NULL prices with 0.	Null handling	Easy	Provide default values for missing data

```
SELECT listing_id, COALESCE(price, 0) AS price  
FROM listings;
```

SQL Topology

0. SQL FUNDAMENTALS

0.1 SQL Syntax and Commands

0.1.1 SELECT and FROM Statements

0.1.2 Simple Filtering with WHERE

0.1.3 Sorting with ORDER BY

0.1.4 Limiting Results with LIMIT/OFFSET

0.1.5 INSERT, UPDATE, and DELETE Statements

0.1.6 TRUNCATE and COPY Statements

0.2 Introduction to Databases

0.2.1 What is SQL and Relational Databases?

0.2.2 Database Management Systems (DBMS) Overview

0.2.3 Query Execution Basics

0.2.4 Creating, Dropping, and Renaming Databases

0.3 Introduction to Tables

0.3.1 Creating, Dropping, Renaming, and Altering Tables

0.3.2 Temporary Tables

0.3.3 Sequences and Auto-Increment

1. Database Architecture and Objects

1.1.1 Schemas, Tables, Views, and Functions

1.1.2 Primary Keys, Foreign Keys, and Relationships

1.1.3 Indexes and Performance Optimization

1.1.4 Stored Procedures, Triggers, and Cursors

1.2 SQL Data Types and Constraints

1.2.1 Numeric, String, Date/Time, and Boolean Types

1.2.2 Data Type Conversions and Casting

1.2.3 Constraints for Data Integrity

1.3 Database Types for Analytics (Row-store vs. Column-store)

2. DATA PREPARATION AND PROFILING

2.1 Exploratory Data Analysis with SQL

2.1.1 Data Profiling Techniques

2.1.2 Frequency Analysis and Distributions

2.1.3 Identifying Data Quality Issues

2.2 Data Cleaning and Transformation

2.2.1 Handling Null Values and Missing Data

2.2.2 String Manipulation and Text Processing

2.2.3 Date and Time Manipulations

2.2.4 Importing and Exporting Data

2.3 Data Shaping and Restructuring

2.3.1 Pivoting and Unpivoting Data

2.3.2 Data Granularity and Aggregation Levels

2.3.3 Creating Analytical Datasets

3. CORE SQL FOR ANALYTICS

3.1 Advanced SELECT Operations

3.1.1 Complex WHERE Clauses and Filtering

3.1.2 Pattern Matching with LIKE and Regular Expressions

3.1.3 Subqueries and Correlated Subqueries

3.1.4 Nested Queries

3.1.5 Operators (IN, BETWEEN, EXISTS, ALL, ANY, NOT, IS NULL, Wildcards)

3.1.6 Set Operations (UNION, UNION ALL, INTERSECT, EXCEPT)

3.2 JOIN Strategies for Analysis

3.2.1 Inner, Left, Right, and Full Outer Joins

- 3.2.2 Self-Joins and Cross Joins
- 3.2.3 Complex Multi-Table Joins
- 3.2.4 Semi-Joins and Anti-Joins
- 3.2.5 Theta Joins, Equi-Joins, Non-Equi-Joins, and Natural Joins
- 3.2.6 Implicit and Explicit Joins
- 3.3 Aggregation and Grouping
 - 3.3.1 GROUP BY and HAVING Clauses
 - 3.3.2 Aggregate Functions (COUNT, SUM, AVG, MIN, MAX)
 - 3.3.3 Grouping Sets, Rollup, and Cube Operations
 - 3.3.4 Median and Percentiles

4. ADVANCED ANALYTICAL FUNCTIONS

- 4.1 Window Functions for Analytics
 - 4.1.1 ROW_NUMBER, RANK, and DENSE_RANK
 - 4.1.2 LAG and LEAD for Time-Based Analysis
 - 4.1.3 Cumulative and Moving Aggregations
 - 4.1.4 NTILE for Bucketing
- 4.2 Common Table Expressions (CTEs)
 - 4.2.1 Basic CTEs for Query Organization
 - 4.2.2 Recursive CTEs
 - 4.2.3 CTEs vs. Subqueries vs. Temp Tables
- 4.3 CASE Statements and Conditional Logic
 - 4.3.1 Simple and Searched CASE Expressions
 - 4.3.2 Data Categorization and Binning
 - 4.3.3 Flag Creation and Binary Encoding
 - 4.3.4 Conditional Functions (IF, COALESCE, NULLIF, DECODE)

5. TIME SERIES ANALYSIS WITH SQL

5.1 Date and Time Functions

5.1.1 Date Arithmetic and Calculations

5.1.2 Date Parts Extraction and Formatting

5.1.3 Time Zone Handling

5.1.4 Date Functions (NOW, CURDATE, CURTIME, DATEPART, EXTRACT, DATE_ADD, DATEDIFF, DATE_FORMAT)

5.2 Temporal Analytics

5.2.1 Period-over-Period Comparisons

5.2.2 Rolling Time Windows and Moving Averages

5.2.3 Seasonality Analysis

5.2.4 Week-over-Week (WoW), Month-over-Month (MoM), and Year-over-Year (YoY) Calculations

5.3 Time Series Data Preparation

5.3.1 Date Dimensions and Calendar Tables

5.3.2 Gap Analysis and Missing Time Periods

5.3.3 Data Granularity Adjustments

5.3.4 Relative Dates and Time Durations

6. COHORT AND RETENTION ANALYSIS

6.1 Cohort Analysis Framework

6.1.1 Defining Cohorts and Time Series

6.1.2 Cohort vs. Segment Distinctions

6.1.3 Aggregate Metrics for Cohorts

6.2 Retention Analytics

6.2.1 Basic Retention Curves

6.2.2 Sparse Data Handling

6.2.3 Retention by Different Time Periods

6.3 Advanced Cohort Metrics

6.3.1 Survivorship Analysis

6.3.2 Returnship and Repeat Behavior

6.3.3 Customer Lifetime Value (CLTV) Calculations

7. BUSINESS INTELLIGENCE AND KPI DEVELOPMENT

7.1 KPI Framework Design

7.1.1 Metric Definitions and Calculations

7.1.2 Performance Indicators and Benchmarks

7.1.3 Dashboard-Ready Data Structures

7.2 Business Metrics and Analytics

7.2.1 Sales and Revenue Analysis

7.2.2 Customer Analytics and Segmentation

7.2.3 Marketing Attribution and Campaign Analysis

7.2.4 Ads Performance and Friend Requests Analysis

7.2.5 Average Duration Between Visits

7.3 Reporting and Visualization Preparation

7.3.1 Data Aggregation for Dashboards

7.3.2 Drill-down and Roll-up Capabilities

7.3.3 Performance Optimization for Reports

8. ADVANCED ANALYTICAL TECHNIQUES

8.1 Statistical Analysis with SQL

8.1.1 Descriptive Statistics and Percentiles

8.1.2 Correlation and Basic Statistical Tests

8.1.3 Outlier Detection and Anomaly Analysis

8.1.4 Ratios and Proportions

8.2 Text Analytics and Pattern Recognition

8.2.1 String Functions and Text Processing

8.2.2 Regular Expressions for Data Extraction

- 8.2.3 Text Classification and Categorization
- 8.3 Experiment Analysis and A/B Testing
 - 8.3.1 Randomized Experiment Design
 - 8.3.2 Control and Treatment Group Analysis
 - 8.3.3 Statistical Significance and Confidence Intervals
- 8.4 Spatial Data Analysis
 - 8.4.1 Introduction to Spatial Data with PostGIS
 - 8.4.2 Working with Spatial Objects (Points, Lines, Polygons)

9. DATA MODELING FOR ANALYTICS

- 9.1 Dimensional Modeling Concepts
 - 9.1.1 Star Schema and Snowflake Schema
 - 9.1.2 Fact Tables and Dimension Tables
 - 9.1.3 Slowly Changing Dimensions
- 9.2 Analytics-Focused Data Models
 - 9.2.1 Granularity Considerations
 - 9.2.2 Aggregation Tables and Summary Views
 - 9.2.3 Data Marts for Specific Business Areas
- 9.3 Data Warehouse Optimization
 - 9.3.1 Indexing Strategies for Analytics
 - 9.3.2 Partitioning and Clustering
 - 9.3.3 Materialized Views and Caching
- 9.4 Relational Modeling Concepts
 - 9.4.1 Entity-Relationship (ER) Model and Diagrams
 - 9.4.2 Database Normalization and Normal Forms
 - 9.4.3 Denormalization Techniques
 - 9.4.4 Functional Dependencies and Inference Rules

10. PERFORMANCE OPTIMIZATION AND SCALABILITY

10.1 Query Performance Tuning

10.1.1 Execution Plan Analysis

10.1.2 Query Optimization Techniques

10.1.3 Index Usage and Design

10.1.4 Reading and Interpreting Execution Plans

10.1.5 Rule-Based vs. Cost-Based Optimizers

10.2 Big Data SQL Techniques

10.2.1 Handling Large Datasets

10.2.2 Partitioning and Sharding Strategies

10.2.3 Parallel Query Execution

10.3 Resource Management

10.3.1 Query Limits and Sampling

10.3.2 Memory and Storage Optimization

10.3.3 ETL vs. Real-time Query Decisions

10.3.4 Database Backup and Maintenance

11. MACHINE LEARNING DATASET PREPARATION

11.1 Feature Engineering with SQL

11.1.1 Creating Model-Ready Datasets

11.1.2 Feature Selection and Transformation

11.1.3 Training and Test Data Splits

11.2 Time Series Modeling Datasets

11.2.1 Forecasting Data Structures

11.2.2 Lag Features and Time Windows

11.2.3 Seasonal and Trend Components

11.3 Classification and Regression Datasets

11.3.1 Binary and Multi-class Targets

11.3.2 Feature Engineering Techniques

11.3.3 Data Preprocessing and Scaling

12. ADVANCED SQL PATTERNS FOR ANALYSTS

12.1 Complex Query Structures

12.1.1 UNION and Set Operations

12.1.2 Self-Joins for Advanced Analysis

12.1.3 Recursive Queries and Hierarchical Data

12.1.4 Dynamic SQL

12.2 Funnel Analysis and User Journey

12.2.1 Multi-step Process Analysis

12.2.2 Conversion Rate Optimization

12.2.3 Drop-off Point Identification

12.2.4 Status Change Analysis

12.3 Basket Analysis and Market Research

12.3.1 Association Rules and Co-occurrence

12.3.2 Product Affinity Analysis

12.3.3 Customer Behavior Patterns

12.4 Top-N Queries and Rankings

13. DATA GOVERNANCE AND QUALITY

13.1 Data Quality Assurance

13.1.1 Data Validation Techniques

13.1.2 Consistency Checks and Constraints

13.1.3 Data Profiling for Quality Assessment

13.2 Privacy and Security Considerations

13.2.1 PII Handling and Anonymization

13.2.2 Data Access Controls and Permissions

13.2.3 Compliance (GDPR, CCPA) Considerations

13.3 Documentation and Metadata Management

13.3.1 Data Dictionary Development

13.3.2 Query Documentation Standards

13.3.3 Version Control for SQL Code

13.3.4 SQL Comments and Documentation

14. TOOL INTEGRATION AND WORKFLOW

14.1 SQL with Business Intelligence Tools

14.1.1 Tableau, Power BI, and Looker Integration

14.1.2 Optimal Data Structures for BI Tools

14.1.3 Performance Considerations

14.2 SQL with Programming Languages

14.2.1 Python Integration (pandas, SQLAlchemy)

14.2.2 R Integration (dplyr, RPostgreSQL)

14.2.3 Jupyter Notebooks and SQL

14.3 Cloud Analytics Platforms

14.3.1 Amazon Redshift and AWS Analytics

14.3.2 Snowflake for Data Warehousing

14.4 Command-Line Tools and Interfaces

14.4.1 Using psql and Other Command-Line Clients

14.4.2 Database Hosting and Connections

15. PRACTICAL APPLICATION AND CASE STUDIES

15.1 Real-World Analytics Projects

15.1.1 Customer Segmentation Analysis

15.1.2 Marketing Campaign Effectiveness

15.1.3 Identifying and Telling the Story Behind Your Data

15.2 Transactions and Concurrency

15.2.1 SQL Transactions

15.2.2 ACID Properties and Management

SQL Theory

SQL THEORY

Core Query Structure

CLAUSE	PURPOSE	INTERVIEW
SELECT	Specifies columns, aggregates, and calculated metrics. The backbone of all data retrieval.	...
FROM	Identifies the data source. Uses the structure: catalog.schema.table (e.g., hive.marketing.campaigns).	...
WHERE	Filters rows before aggregation or grouping.	...
GROUP BY	Aggregates rows based on non-aggregated columns.	Used for standard reporting metrics (sum of sales per region). Can be extended with ROLLUP or CUBE for multi-level aggregates.
HAVING	Filters the results after aggregation/grouping.	Used to filter groups (e.g., HAVING COUNT(user_id) > 100) Rule: WHERE filters rows, HAVING filters groups.
WITH (CTE)	Creates Common Table Expressions (CTEs) to name and define subqueries.	Mandatory for complex queries Improves readability and allows you to break down logic into reusable steps

LIMIT / OFFSET	Restricts the number of output rows.	Used for efficient testing, prototyping, and implementing pagination logic in applications.
---------------------------	--------------------------------------	---

Analytical Window Functions

FUNCTION	CATEGORY	EXAMPLE	ANALYTICS USE CASE
<code>ROW_NUMBER()</code>	Ranking	<code>ROW_NUMBER() OVER (PARTITION BY user_id ORDER BY transaction_date DESC)</code>	Finding the most recent (or first) record per group. No ties possible.
<code>RANK()</code>	Ranking	<code>RANK() OVER (PARTITION BY dept ORDER BY salary DESC)</code>	Assigns a rank. Ties get the same rank, and the next rank number is skipped (e.g., 1, 2, 2, 4).
<code>DENSE_RANK()</code>	Ranking	<code>DENSE_RANK() OVER (PARTITION BY dept ORDER BY salary DESC)</code>	Assigns a rank. Ties get the same rank, and the next rank number is not skipped (e.g., 1, 2, 2, 3).
<code>LAG()</code>	Value	<code>LAG(sales_amount, 1) OVER (ORDER BY date)</code>	Calculating period-over-period change, such as sales from the previous day, month, or quarter.
<code>LEAD()</code>	Value	<code>LEAD(price, 1) OVER (PARTITION BY stock_id ORDER BY timestamp)</code>	Looking ahead to the next value in a sequence (e.g., finding the next price increase).
<code>SUM() / AVG()</code>	Aggregate	<code>SUM(sales) OVER (ORDER BY week ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW)</code>	Calculating running totals or cumulative sums over time.

Join Types

FUNCTION	CATEGORY	EXAMPLE
INNER JOIN	Returns only rows where the join key exists in both tables	Used for merging core data tables where a match is necessary. Generally the fastest join type
LEFT JOIN	Returns all rows from the left table and matched rows from the right table (filling in NULL for unmatched rows)	Used for data enrichment, such as joining a master list of products to a sparse table of reviews
FULL OUTER JOIN	Returns all rows from both tables, filling in NULLs for any non-matches	Used to find and analyze data discrepancies or missing records between two datasets
CROSS JOIN	Returns the Cartesian product (every row from the first table matched with every row from the second table)	Extremely rare in analytics; typically only used for generating scaffolding data (e.g., a complete date-product pair combination)
...		

...
-- Single condition
SELECT * FROM bookings

```
WHERE booking_date >= DATE '2024-01-01';

-- Multiple conditions
SELECT * FROM bookings
WHERE status = 'confirmed'
    AND total_price > 500
    AND guest_country = 'US';

-- IN operator
SELECT * FROM listings
WHERE city IN ('San Francisco', 'New York', 'Los Angeles');

-- LIKE pattern matching
SELECT * FROM listings
WHERE property_type LIKE '%Apartment%';

-- IS NULL / IS NOT NULL
SELECT * FROM reviews
WHERE response_date IS NULL;
```

Comprehensive SQL Guide for Analytics

SQL FUNDAMENTALS

SQL Syntax and Commands

SELECT and FROM Statements

SELECT retrieves data from tables, FROM specifies the source table(s).

```
SELECT column1, column2, ...
FROM table_name;
```

0.1.2 Simple Filtering with WHERE

Purpose: Filter rows based on conditions before any grouping occurs.

-- Basic WHERE clause

```
SELECT * FROM customers
```

```
WHERE country = 'Germany';
```

```
-- Multiple conditions with operators  
  
SELECT product_name, price  
  
FROM products  
  
WHERE price > 10 AND category_id = 2;
```

```
-- Pattern matching with LIKE  
  
SELECT * FROM customers  
  
WHERE customer_name LIKE 'A%'; -- Starts with 'A'
```

Important Notes:

- Presto uses single quotes for string literals
- String comparisons are case-sensitive in Presto
- WHERE filters individual rows before GROUP BY

0.1.3 Sorting with ORDER BY

```
-- Single column sorting  
  
SELECT * FROM products  
  
ORDER BY price DESC;
```

```
-- Multiple column sorting  
  
SELECT customer_name, country, city  
  
FROM customers  
  
ORDER BY country ASC, city DESC;
```

```
-- Presto supports NULLS FIRST/LAST  
  
SELECT * FROM products  
  
ORDER BY discount DESC NULLS LAST;
```

0.1.4 Limiting Results with LIMIT/OFFSET

-- Get top 10 results (Presto syntax)

```
SELECT * FROM orders
```

```
LIMIT 10;
```

-- Pagination: Skip 20, get next 10

```
SELECT * FROM orders
```

```
LIMIT 10 OFFSET 20;
```

-- Alternative syntax

```
SELECT * FROM orders
```

```
OFFSET 20 LIMIT 10;
```

0.1.5 INSERT, UPDATE, and DELETE Statements

Note: Presto is primarily for analytics (read-only). These operations are typically done in Hive or the underlying data store.

-- INSERT: Add new data (Hive syntax)

```
INSERT INTO customers VALUES ('Cardinal', 'Tom B. Erichsen', 'Norway');
```

-- INSERT with specific columns

```
INSERT INTO customers (customer_name, contact_name, country)
```

```
VALUES ('Cardinal', 'Tom B. Erichsen', 'Norway');
```

-- UPDATE and DELETE are not directly supported in Presto

```
-- Use Hive or recreate tables with CREATE TABLE AS
```

0.1.6 TRUNCATE and COPY Statements

```
-- In Presto, recreate table instead of TRUNCATE
```

```
DROP TABLE IF EXISTS temp_customers;
```

```
CREATE TABLE temp_customers AS SELECT * FROM customers WHERE 1=0;
```

```
-- Copy data between tables
```

```
CREATE TABLE new_table AS
```

```
SELECT * FROM old_table;
```

0.2 Introduction to Databases

0.2.1 What is SQL and Relational Databases?

- **SQL:** Structured Query Language for managing relational databases
- **Relational Database:** Stores data in structured tables with relationships
- **Key Benefits:**
 - Widely used across all organizations
 - Skills transferable between database systems
 - ACID properties ensure data consistency

0.2.2 Database Management Systems (DBMS) Overview

- **Presto/Trino:** Distributed SQL query engine for big data
- **Hive:** Data warehouse infrastructure built on Hadoop
- **OLTP vs OLAP:**
 - **OLTP:** High volume of short transactions, normalized schemas
 - **OLAP:** Complex queries on large datasets, denormalized schemas (Presto's domain)

0.2.3 Query Execution Basics

Logical Execution Order (Critical for understanding):

1. FROM / JOIN - Identifies and combines source tables
2. WHERE - Filters individual rows
3. GROUP BY - Aggregates rows into groups
4. HAVING - Filters aggregated groups
5. SELECT - Selects final columns and calculations
6. DISTINCT - Removes duplicates
7. ORDER BY - Sorts result set
8. LIMIT/OFFSET - Restricts rows returned

0.2.4 Creating, Dropping, and Renaming Databases

-- In Presto, work with schemas instead of databases

```
CREATE SCHEMA IF NOT EXISTS analytics;
```

```
DROP SCHEMA IF EXISTS test_schema;
```

-- Switch schema

```
USE analytics;
```

0.3 Introduction to Tables

0.3.1 Creating, Dropping, Renaming, and Altering Tables

-- Create table (Hive/Presto syntax)

```
CREATE TABLE IF NOT EXISTS employees (
```

```
    employee_id BIGINT,
```

```
    name VARCHAR,
```

```
    email VARCHAR,
```

```
    salary DECIMAL(10,2),
```

```
    hire_date DATE
```

```
) WITH (
```

```
    format = 'ORC',
```

```
partitioned_by = ARRAY['hire_date']

);

-- Drop table

DROP TABLE IF EXISTS temp_employees;

-- Create table as select (CTAS)

CREATE TABLE high_value_customers AS

SELECT customer_id, SUM(amount) as total_spent

FROM orders

GROUP BY customer_id

HAVING SUM(amount) > 10000;
```

```
-- Presto doesn't support ALTER TABLE for structure changes

-- Recreate table with new structure instead
```

0.3.2 Temporary Tables

```
-- Presto uses WITH clauses (CTEs) instead of temp tables

WITH temp_calculations AS (

    SELECT customer_id, SUM(amount) as total_spent

    FROM orders

    GROUP BY customer_id

)

SELECT * FROM temp_calculations;
```

0.3.3 Sequences and Auto-Increment

```
-- Presto doesn't have auto-increment  
-- Generate sequences using ROW_NUMBER()  
  
SELECT  
    ROW_NUMBER() OVER (ORDER BY created_at) as id,  
    product_name  
FROM products;
```

1. DATABASE ARCHITECTURE AND OBJECTS

1.1 Database Structure

1.1.1 Schemas, Tables, Views, and Functions

```
-- Create schema  
CREATE SCHEMA IF NOT EXISTS marketing;  
  
-- Create view  
CREATE OR REPLACE VIEW customer_orders AS  
SELECT c.customer_name, o.order_date, o.total_amount  
FROM customers c  
JOIN orders o ON c.customer_id = o.customer_id;
```

```
-- Presto doesn't support stored functions  
-- Use inline calculations or CTEs instead
```

1.1.2 Primary Keys, Foreign Keys, and Relationships

```
-- Presto doesn't enforce constraints
```

```
-- Document relationships in table comments  
  
COMMENT ON TABLE orders IS 'Orders table with customer_id as FK to  
customers.customer_id';
```

1.1.3 Indexes and Performance Optimization

```
-- Presto doesn't create indexes  
  
-- Performance optimization through:  
  
-- 1. Partitioning  
  
-- 2. Bucketing  
  
-- 3. File formats (ORC, Parquet)  
  
-- 4. Statistics
```

```
-- Create partitioned table
```

```
CREATE TABLE events (  
    event_id BIGINT,  
    user_id BIGINT,  
    event_type VARCHAR,  
    event_time TIMESTAMP  
) WITH (  
    format = 'ORC',  
    partitioned_by = ARRAY['event_date'],  
    bucketed_by = ARRAY['user_id'],  
    bucket_count = 50  
);
```

1.2 SQL Data Types and Constraints

1.2.1 Numeric, String, Date/Time, and Boolean Types

Presto Data Types:

- **Numeric:** TINYINT, SMALLINT, INTEGER, BIGINT, REAL, DOUBLE, DECIMAL(p,s)
- **String:** VARCHAR, CHAR(n), VARBINARY, JSON
- **Date/Time:** DATE, TIME, TIMESTAMP, INTERVAL
- **Boolean:** BOOLEAN
- **Complex:** ARRAY, MAP, ROW

1.2.2 Data Type Conversions and Casting

-- Explicit casting in Presto

SELECT

```
CAST('123' AS INTEGER) as int_value,  
CAST(order_date AS VARCHAR) as date_string,  
CAST(price AS DECIMAL(10,2)) as formatted_price,  
TRY_CAST('invalid' AS INTEGER) as safe_cast -- Returns NULL instead of error  
FROM orders;
```

2. DATA PREPARATION AND PROFILING

2.1 Exploratory Data Analysis with SQL

2.1.1 Data Profiling Techniques

-- Basic profiling query

SELECT

```
COUNT(*) as total_rows,  
COUNT(DISTINCT customer_id) as unique_customers,  
MIN(order_date) as earliest_order,
```

```
MAX(order_date) as latest_order,  
AVG(total_amount) as avg_order_value,  
APPROX_PERCENTILE(total_amount, 0.5) as median_order_value  
FROM orders;
```

```
-- Check for nulls  
SELECT  
    COUNT(*) - COUNT(customer_id) as null_customers,  
    COUNT(*) - COUNT(order_date) as null_dates  
FROM orders;
```

2.1.2 Frequency Analysis and Distributions

```
-- Frequency distribution  
SELECT  
    status,  
    COUNT(*) as frequency,  
    COUNT(*) * 100.0 / SUM(COUNT(*)) OVER () as percentage  
FROM orders  
GROUP BY status  
ORDER BY frequency DESC;
```

```
-- Histogram bins using WIDTH_BUCKET  
SELECT  
    WIDTH_BUCKET(price, 0, 100, 10) as price_bucket,  
    COUNT(*) as count
```

```
FROM products  
GROUP BY 1  
ORDER BY 1;
```

2.1.3 Identifying Data Quality Issues

```
-- Find duplicates  
  
SELECT customer_email, COUNT(*) as count  
  
FROM customers  
  
GROUP BY customer_email  
  
HAVING COUNT(*) > 1;
```

```
-- Find orphaned records (anti-join pattern)  
  
SELECT o.*  
  
FROM orders o  
  
LEFT JOIN customers c ON o.customer_id = c.customer_id  
  
WHERE c.customer_id IS NULL;
```

2.2 Data Cleaning and Transformation

2.2.1 Handling Null Values and Missing Data

```
-- Replace nulls with default value using COALESCE  
  
SELECT  
  
    customer_id,  
  
    COALESCE(phone, 'No Phone') as phone,  
  
    COALESCE(email, 'No Email') as email,  
  
    IF(address IS NULL, 'No Address', address) as address
```

```
FROM customers;

-- Filter out nulls

SELECT * FROM orders

WHERE ship_date IS NOT NULL;

-- Use NULLIF to handle specific values as NULL

SELECT

    revenue,

    expenses,

    revenue / NULLIF(expenses, 0) as profit_ratio -- Prevents division by zero

FROM financial_data;
```

2.2.2 String Manipulation and Text Processing

```
-- Common string functions in Presto

SELECT

    UPPER(customer_name) as uppercase_name,

    LOWER(email) as lowercase_email,

    TRIM(address) as trimmed_address,

    SUBSTR(phone, 1, 3) as area_code,

    CONCAT(first_name, ' ', last_name) as full_name,

    first_name || ' ' || last_name as full_name_alt,

    REPLACE(description, 'old', 'new') as updated_desc,

    SPLIT(tags, ',') as tag_array,

    LENGTH(description) as desc_length
```

```

FROM customers;

-- Regular expressions in Presto

SELECT * FROM customers

WHERE REGEXP_LIKE(email, '^([A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Z|a-z]{2,})$');

-- Extract with regex

SELECT

    REGEXP_EXTRACT(url, 'utm_source=([^\&]+)') as utm_source,
    REGEXP_EXTRACT_ALL(text, '\d+') as all_numbers

FROM marketing_data;

```

2.2.3 Date and Time Manipulations

```

-- Date functions in Presto

SELECT

    CURRENT_DATE as today,
    CURRENT_TIMESTAMP as right_now,
    DATE_ADD('day', 7, order_date) as week_later,
    DATE_DIFF('day', order_date, ship_date) as days_to_ship,
    EXTRACT(YEAR FROM order_date) as order_year,
    EXTRACT(MONTH FROM order_date) as order_month,
    DATE_TRUNC('month', order_date) as month_start,
    DATE_FORMAT(order_date, '%Y-%m') as year_month,
    order_date + INTERVAL '30' DAY as payment_due

FROM orders;

```

```
-- Working with timestamps

SELECT

    FROM_UNIXTIME(unix_timestamp) as timestamp_value,

    TO_UNIXTIME(timestamp_col) as unix_value

FROM events;
```

2.3 Data Shaping and Restructuring

2.3.1 Pivoting and Unpivoting Data

```
-- Pivot using conditional aggregation (Presto doesn't have PIVOT)

SELECT

    customer_id,

    SUM(CASE WHEN YEAR(order_date) = 2023 THEN revenue END) as revenue_2023,
    SUM(CASE WHEN YEAR(order_date) = 2024 THEN revenue END) as revenue_2024,
    SUM(CASE WHEN YEAR(order_date) = 2025 THEN revenue END) as revenue_2025

FROM sales

GROUP BY customer_id;
```

```
-- Unpivot using CROSS JOIN with VALUES

SELECT customer_id, year, revenue

FROM sales

CROSS JOIN UNNEST(
    ARRAY[2023, 2024, 2025],
    ARRAY[revenue_2023, revenue_2024, revenue_2025]
) AS t(year, revenue);
```

3. CORE SQL FOR ANALYTICS

3.1 Advanced SELECT Operations

3.1.1 Complex WHERE Clauses and Filtering

-- Complex conditions

SELECT * FROM orders

WHERE (status = 'Shipped' OR status = 'Delivered')

AND order_date >= DATE '2024-01-01'

AND total_amount > 100

AND customer_id IN (

SELECT customer_id FROM customers

WHERE country = 'USA'

);

3.1.2 Pattern Matching with LIKE and Regular Expressions

-- LIKE patterns

SELECT * FROM products

WHERE product_name LIKE 'Ch%' -- Starts with Ch

OR product_name LIKE '%ese%' -- Ends with ese

OR product_name LIKE '%ee%'; -- Contains ee

-- Regular expressions in Presto

SELECT * FROM emails

WHERE REGEXP_LIKE(email_address, '^[a-zA-Z0-9._%+-]+@[airbnb\\.com]\$');

3.1.3 Subqueries and Correlated Subqueries

-- Non-correlated subquery

```
SELECT * FROM employees
```

```
WHERE salary > (SELECT AVG(salary) FROM employees);
```

-- Correlated subquery

```
SELECT e1.name, e1.salary
```

```
FROM employees e1
```

```
WHERE e1.salary > (
```

```
    SELECT AVG(e2.salary)
```

```
    FROM employees e2
```

```
    WHERE e2.department = e1.department
```

```
);
```

3.1.5 Operators (IN, BETWEEN, EXISTS, ALL, ANY, NOT, IS NULL)

-- Various operators in Presto

```
SELECT * FROM products
```

```
WHERE price BETWEEN 10 AND 50
```

```
AND category_id IN (1, 2, 3)
```

```
AND EXISTS (
```

```
    SELECT 1 FROM order_details
```

```
    WHERE order_details.product_id = products.product_id
```

```
)
```

```
AND price > ALL (SELECT price FROM products WHERE category_id = 4)
```

AND discontinued IS NOT NULL;

-- ANY/SOME operator

SELECT * FROM products

WHERE price > ANY (SELECT price FROM products WHERE category_id = 2);

3.1.6 Set Operations (UNION, UNION ALL, INTERSECT, EXCEPT)

-- UNION removes duplicates

SELECT country FROM customers

UNION

SELECT country FROM suppliers;

-- UNION ALL keeps duplicates (faster)

SELECT city FROM customers

UNION ALL

SELECT city FROM suppliers;

-- INTERSECT: rows in both

SELECT product_id FROM orders_2023

INTERSECT

SELECT product_id FROM orders_2024;

-- EXCEPT: rows in first but not second

SELECT customer_id FROM all_customers

EXCEPT

```
SELECT customer_id FROM inactive_customers;
```

3.2 JOIN Strategies for Analysis

3.2.1 Inner, Left, Right, and Full Outer Joins

-- INNER JOIN: Only matching records

```
SELECT c.customer_name, o.order_date  
FROM customers c  
INNER JOIN orders o ON c.customer_id = o.customer_id;
```

-- LEFT JOIN: All from left, matching from right

```
SELECT c.customer_name, o.order_date  
FROM customers c  
LEFT JOIN orders o ON c.customer_id = o.customer_id;
```

-- RIGHT JOIN: All from right, matching from left

```
SELECT c.customer_name, o.order_date  
FROM customers c  
RIGHT JOIN orders o ON c.customer_id = o.customer_id;
```

-- FULL OUTER JOIN: All from both

```
SELECT c.customer_name, o.order_date  
FROM customers c  
FULL OUTER JOIN orders o ON c.customer_id = o.customer_id;
```

3.2.2 Self-Joins and Cross Joins

```
-- Self-join: Find employees with same salary

SELECT e1.name as employee1, e2.name as employee2, e1.salary
FROM employees e1
JOIN employees e2 ON e1.salary = e2.salary
WHERE e1.employee_id < e2.employee_id;
```

```
-- Cross join: Cartesian product

SELECT p.product_name, c.category_name
FROM products p
CROSS JOIN categories c;
```

```
-- Alternative syntax

SELECT p.product_name, c.category_name
FROM products p, categories c;
```

3.2.3 Complex Multi-Table Joins

```
-- Join multiple tables

SELECT
    c.customer_name,
    o.order_date,
    od.quantity,
    p.product_name,
    s.company_name as supplier
FROM customers c
JOIN orders o ON c.customer_id = o.customer_id
```

```
JOIN order_details od ON o.order_id = od.order_id
JOIN products p ON od.product_id = p.product_id
JOIN suppliers s ON p.supplier_id = s.supplier_id
WHERE o.order_date >= DATE '2024-01-01';
```

3.2.4 Semi-Joins and Anti-Joins

```
-- Semi-join using EXISTS
```

```
SELECT * FROM customers c
WHERE EXISTS (
    SELECT 1 FROM orders o
    WHERE o.customer_id = c.customer_id
);
```

```
-- Semi-join using IN
```

```
SELECT * FROM customers
WHERE customer_id IN (SELECT customer_id FROM orders);
```

```
-- Anti-join: Customers with no orders
```

```
SELECT c.*
FROM customers c
LEFT JOIN orders o ON c.customer_id = o.customer_id
WHERE o.customer_id IS NULL;
```

```
-- Anti-join using NOT EXISTS
```

```
SELECT * FROM customers c
```

```
WHERE NOT EXISTS (
    SELECT 1 FROM orders o
    WHERE o.customer_id = c.customer_id
);
```

3.3 Aggregation and Grouping

3.3.1 GROUP BY and HAVING Clauses

```
-- GROUP BY with HAVING

SELECT
    department,
    COUNT(*) as employee_count,
    AVG(salary) as avg_salary
FROM employees
WHERE status = 'Active'
GROUP BY department
HAVING COUNT(*) > 5
    AND AVG(salary) > 50000
ORDER BY avg_salary DESC;
```

Key Difference: WHERE filters rows before grouping, HAVING filters groups after aggregation

3.3.2 Aggregate Functions (COUNT, SUM, AVG, MIN, MAX)

```
SELECT
    COUNT(*) as total_orders,
    COUNT(DISTINCT customer_id) as unique_customers,
    SUM(total_amount) as total_revenue,
```

```
    AVG(total_amount) as avg_order_value,  
    MIN(order_date) as first_order,  
    MAX(order_date) as last_order,  
    ARBITRARY(status) as sample_status, -- Presto-specific  
    BOOL_AND(is_shipped) as all_shipped, -- Presto-specific  
    BOOL_OR(is_expedited) as any_expedited -- Presto-specific  
FROM orders  
WHERE status = 'Completed';
```

3.3.3 Grouping Sets, Rollup, and Cube Operations

```
-- GROUPING SETS in Presto  
  
SELECT  
    product,  
    region,  
    year,  
    SUM(sales) as total,  
    GROUPING(product, region, year) as grouping_id  
  
FROM sales_data  
  
GROUP BY GROUPING SETS (  
    (product, region),  
    (product, year),  
    (region, year),  
    () -- Grand total  
);
```

```
-- ROLLUP  
SELECT  
    year,  
    quarter,  
    month,  
    SUM(revenue) as total  
FROM sales  
GROUP BY ROLLUP (year, quarter, month);
```

```
-- CUBE  
SELECT  
    product,  
    region,  
    SUM(sales) as total  
FROM sales_data  
GROUP BY CUBE (product, region);
```

3.3.4 Median and Percentiles

```
-- Percentiles in Presto  
SELECT  
    APPROX_PERCENTILE(salary, 0.5) as median_salary,  
    APPROX_PERCENTILE(salary, 0.25) as q1,  
    APPROX_PERCENTILE(salary, 0.75) as q3,  
    APPROX_PERCENTILE(salary, ARRAY[0.25, 0.5, 0.75, 0.95]) as percentiles  
FROM employees;
```

4. ADVANCED ANALYTICAL FUNCTIONS

4.1 Window Functions for Analytics

4.1.1 ROW_NUMBER, RANK, and DENSE_RANK

-- Ranking functions in Presto

```
SELECT  
    department,  
    employee_name,  
    salary,  
    ROW_NUMBER() OVER (PARTITION BY department ORDER BY salary DESC) as row_num,  
    RANK() OVER (PARTITION BY department ORDER BY salary DESC) as rank_num,  
    DENSE_RANK() OVER (PARTITION BY department ORDER BY salary DESC) as dense_rank,  
    PERCENT_RANK() OVER (PARTITION BY department ORDER BY salary DESC) as  
    percent_rank  
FROM employees;
```

4.1.2 LAG and LEAD for Time-Based Analysis

-- Period-over-period comparison

```
SELECT  
    month,  
    revenue,  
    LAG(revenue, 1) OVER (ORDER BY month) as prev_month,  
    LAG(revenue, 12) OVER (ORDER BY month) as same_month_last_year,  
    LEAD(revenue, 1) OVER (ORDER BY month) as next_month,  
    revenue - LAG(revenue, 1) OVER (ORDER BY month) as mom_change,
```

```

(revenue - LAG(revenue, 1) OVER (ORDER BY month)) * 100.0 /
NULLIF(LAG(revenue, 1) OVER (ORDER BY month), 0) as mom_growth_pct
FROM monthly_revenue;

-- Sessionization with LAG
SELECT
    user_id,
    event_time,
    LAG(event_time) OVER (PARTITION BY user_id ORDER BY event_time) as prev_event,
    CASE
        WHEN DATE_DIFF('minute',
            LAG(event_time) OVER (PARTITION BY user_id ORDER BY event_time),
            event_time) > 30
        THEN 1
        ELSE 0
    END as new_session
FROM user_events;

```

4.1.3 Cumulative and Moving Aggregations

```

-- Running total in Presto
SELECT
    order_date,
    revenue,
    SUM(revenue) OVER (
        ORDER BY order_date

```

```

ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW

) as running_total,
SUM(revenue) OVER (
    ORDER BY order_date
    RANGE BETWEEN INTERVAL '30' DAY PRECEDING AND CURRENT ROW

) as rolling_30_day

FROM daily_revenue;

-- Moving average

SELECT
    date,
    sales,
    AVG(sales) OVER (
        ORDER BY date
        ROWS BETWEEN 6 PRECEDING AND CURRENT ROW
    ) as moving_avg_7d,
    AVG(sales) OVER (
        ORDER BY date
        ROWS BETWEEN 29 PRECEDING AND CURRENT ROW
    ) as moving_avg_30d

FROM daily_sales;

4.1.4 NTILE for Bucketing

-- Divide customers into quartiles

SELECT

```

```
customer_id,  
total_purchases,  
NTILE(4) OVER (ORDER BY total_purchases DESC) as quartile,  
NTILE(10) OVER (ORDER BY total_purchases DESC) as decile,  
NTILE(100) OVER (ORDER BY total_purchases DESC) as percentile  
FROM customer_spending;
```

4.2 Common Table Expressions (CTEs)

4.2.1 Basic CTEs for Query Organization

```
-- Simple CTE  
  
WITH high_value_customers AS (  
    SELECT customer_id, SUM(total_amount) as total_spent  
    FROM orders  
    GROUP BY customer_id  
    HAVING SUM(total_amount) > 10000  
)  
  
SELECT c.customer_name, h.total_spent  
FROM high_value_customers h  
JOIN customers c ON h.customer_id = c.customer_id;
```

```
-- Multiple CTEs  
  
WITH  
customer_totals AS (  
    SELECT customer_id, SUM(amount) as total  
    FROM orders
```

```

        GROUP BY customer_id
    ),
customer_ranks AS (
    SELECT
        customer_id,
        total,
        DENSE_RANK() OVER (ORDER BY total DESC) as rank
    FROM customer_totals
)
SELECT * FROM customer_ranks WHERE rank <= 10;

```

4.2.2 Recursive CTEs

```

-- Note: Presto has limited recursive CTE support

-- Alternative: Use iterative approach or external processing

-- Example of hierarchical data traversal without recursion

WITH employee_hierarchy AS (
    SELECT
        e1.employee_id,
        e1.name,
        e1.manager_id,
        e2.name as manager_name,
        e3.name as skip_manager_name
    FROM employees e1
    LEFT JOIN employees e2 ON e1.manager_id = e2.employee_id
    LEFT JOIN employees e3 ON e2.manager_id = e3.employee_id

```

```
)
```

```
SELECT * FROM employee_hierarchy;
```

4.3 CASE Statements and Conditional Logic

4.3.1 Simple and Searched CASE Expressions

```
-- Simple CASE
```

```
SELECT
```

```
    order_id,
```

```
    status,
```

```
    CASE status
```

```
        WHEN 'Pending' THEN 'Awaiting Processing'
```

```
        WHEN 'Shipped' THEN 'In Transit'
```

```
        WHEN 'Delivered' THEN 'Completed'
```

```
        ELSE 'Unknown'
```

```
    END as status_description
```

```
FROM orders;
```

```
-- Searched CASE
```

```
SELECT
```

```
    customer_id,
```

```
    total_purchases,
```

```
    CASE
```

```
        WHEN total_purchases >= 10000 THEN 'Platinum'
```

```
        WHEN total_purchases >= 5000 THEN 'Gold'
```

```
        WHEN total_purchases >= 1000 THEN 'Silver'
```

```
    ELSE 'Bronze'  
END as customer_tier  
  
FROM customer_summary;
```

4.3.2 Data Categorization and Binning

```
-- Age group binning  
  
SELECT  
  
customer_id,  
  
age,  
  
CASE  
  
WHEN age < 18 THEN 'Under 18'  
WHEN age BETWEEN 18 AND 24 THEN '18-24'  
WHEN age BETWEEN 25 AND 34 THEN '25-34'  
WHEN age BETWEEN 35 AND 44 THEN '35-44'  
WHEN age BETWEEN 45 AND 54 THEN '45-54'  
WHEN age >= 55 THEN '55+'  
  
END as age_group  
  
FROM customers;
```

```
-- Using WIDTH_BUCKET for automatic binning  
  
SELECT  
  
customer_id,  
  
purchase_amount,  
  
WIDTH_BUCKET(purchase_amount, 0, 1000, 10) as amount_bucket  
  
FROM purchases;
```

4.3.4 Conditional Functions (IF, COALESCE, NULLIF)

-- COALESCE: First non-null value

SELECT

```
customer_id,  
COALESCE(email, phone, 'No Contact') as primary_contact  
FROM customers;
```

-- NULLIF: Returns NULL if expressions equal

SELECT

```
revenue,  
expenses,  
revenue / NULLIF(expenses, 0) as profit_ratio -- Prevents division by zero  
FROM financial_data;
```

-- IF function in Presto

SELECT

```
customer_id,  
IF(total_purchases > 5000, 'VIP', 'Regular') as customer_type,  
IF(age >= 18, 'Adult', 'Minor') as age_category  
FROM customer_summary;
```

-- TRY for safe operations

SELECT

```
TRY(CAST(string_col AS INTEGER)) as safe_int,
```

```
TRY(date_col / amount) as safe_division  
FROM data_table;
```

5. TIME SERIES ANALYSIS WITH SQL

5.1 Date and Time Functions

5.1.1 Date Arithmetic and Calculations

```
-- Date arithmetic in Presto  
  
SELECT  
  
    order_date,  
  
    ship_date,  
  
    DATE_DIFF('day', order_date, ship_date) as days_to_ship,  
  
    order_date + INTERVAL '30' DAY as payment_due,  
  
    order_date - INTERVAL '1' YEAR as year_ago,  
  
    DATE_ADD('month', 3, order_date) as quarter_later,  
  
    DATE_ADD('week', -2, order_date) as two_weeks_before  
  
FROM orders;
```

5.1.2 Date Parts Extraction and Formatting

```
-- Extract date parts in Presto  
  
SELECT  
  
    order_date,  
  
    YEAR(order_date) as year,  
  
    MONTH(order_date) as month,  
  
    DAY(order_date) as day,
```

```
DAY_OF_WEEK(order_date) as day_of_week, -- 1=Monday, 7=Sunday  
DAY_OF_YEAR(order_date) as day_of_year,  
WEEK(order_date) as week_number,  
QUARTER(order_date) as quarter,  
DATE_FORMAT(order_date, '%Y-%m') as year_month,  
DATE_FORMAT(order_date, '%Y-W%v') as year_week_iso,  
DATE_TRUNC('month', order_date) as month_start,  
DATE_TRUNC('week', order_date) as week_start  
  
FROM orders;
```

5.1.3 Time Zone Handling

```
-- Time zone operations in Presto  
  
SELECT  
  
    AT_TIMEZONE(event_timestamp, 'America/Los_Angeles') as la_time,  
    AT_TIMEZONE(event_timestamp, 'UTC') as utc_time,  
    WITH_TIMEZONE(event_timestamp, 'America/New_York') as ny_time,  
    CURRENT_TIMEZONE() as session_timezone  
  
FROM events;
```

5.2 Temporal Analytics

5.2.1 Period-over-Period Comparisons

```
-- Year-over-Year comparison  
  
WITH current_year AS (  
    SELECT  
        DATE_TRUNC('month', order_date) as month,
```

```

        SUM(revenue) as revenue
    FROM orders
    WHERE YEAR(order_date) = 2024
    GROUP BY 1
),
previous_year AS (
    SELECT
        DATE_TRUNC('month', order_date + INTERVAL '1' YEAR) as month,
        SUM(revenue) as revenue
    FROM orders
    WHERE YEAR(order_date) = 2023
    GROUP BY 1
)
SELECT
    c.month,
    c.revenue as current_revenue,
    p.revenue as previous_revenue,
    (c.revenue - p.revenue) * 100.0 / NULLIF(p.revenue, 0) as yoy_growth
FROM current_year c
LEFT JOIN previous_year p ON c.month = p.month
ORDER BY c.month;

```

5.2.2 Rolling Time Windows and Moving Averages

-- 30-day rolling revenue

```
SELECT
```

```
date,  
daily_revenue,  
SUM(daily_revenue) OVER (  
    ORDER BY date  
    RANGE BETWEEN INTERVAL '29' DAY PRECEDING AND CURRENT ROW  
) as rolling_30_day  
FROM daily_revenue;
```

-- Week-over-week with proper date handling

```
WITH weekly_metrics AS (
```

```
    SELECT  
        DATE_TRUNC('week', date) as week,  
        SUM(revenue) as weekly_revenue  
    FROM daily_revenue  
    GROUP BY 1
```

```
)
```

```
SELECT
```

```
    week,  
    weekly_revenue,  
    LAG(weekly_revenue, 1) OVER (ORDER BY week) as prev_week,  
    (weekly_revenue - LAG(weekly_revenue, 1) OVER (ORDER BY week)) * 100.0 /  
    NULLIF(LAG(weekly_revenue, 1) OVER (ORDER BY week), 0) as wow_growth  
FROM weekly_metrics;
```

6. COHORT AND RETENTION ANALYSIS

6.1 Cohort Analysis Framework

6.1.1 Defining Cohorts and Time Series

```
-- Define user cohorts by signup month

WITH user_cohorts AS (
    SELECT
        user_id,
        DATE_TRUNC('month', MIN(signup_date)) as cohort_month
    FROM users
    GROUP BY user_id
)

SELECT
    cohort_month,
    COUNT(DISTINCT user_id) as cohort_size
FROM user_cohorts
GROUP BY cohort_month
ORDER BY cohort_month;
```

6.2 Retention Analytics

6.2.1 Basic Retention Curves

```
-- Classic retention analysis in Presto

WITH cohorts AS (
    SELECT
        user_id,
        DATE_TRUNC('month', MIN(event_date)) as cohort_month
    FROM user_events
)
```

```
        GROUP BY user_id
),
activities AS (
    SELECT
        c.cohort_month,
        DATE_DIFF('month', c.cohort_month, DATE_TRUNC('month', e.event_date)) as
months_since,
        COUNT(DISTINCT e.user_id) as active_users
    FROM user_events e
    JOIN cohorts c ON e.user_id = c.user_id
    GROUP BY 1, 2
),
cohort_sizes AS (
    SELECT
        cohort_month,
        COUNT(DISTINCT user_id) as cohort_size
    FROM cohorts
    GROUP BY cohort_month
)
SELECT
    a.cohort_month,
    a.months_since,
    a.active_users,
    cs.cohort_size,
    a.active_users * 100.0 / cs.cohort_size as retention_rate
FROM activities a
```

```
JOIN cohort_sizes cs ON a.cohort_month = cs.cohort_month  
ORDER BY a.cohort_month, a.months_since;
```

6.3 Advanced Cohort Metrics

6.3.3 Customer Lifetime Value (CLTV) Calculations

```
-- Simple LTV calculation  
  
WITH customer_revenue AS (  
  
    SELECT  
  
        customer_id,  
  
        SUM(revenue) as total_revenue,  
  
        COUNT(DISTINCT DATE_TRUNC('month', order_date)) as active_months,  
  
        MIN(order_date) as first_order,  
  
        MAX(order_date) as last_order,  
  
        DATE_DIFF('day', MIN(order_date), MAX(order_date)) as customer_lifespan_days  
  
    FROM orders  
  
    GROUP BY customer_id  
  
)  
  
SELECT  
  
    AVG(total_revenue) as avg_ltv,  
  
    AVG(total_revenue / NULLIF(active_months, 0)) as avg_monthly_value,  
  
    AVG(customer_lifespan_days) as avg_customer_lifespan,  
  
    APPROX_PERCENTILE(total_revenue, 0.5) as median_ltv  
  
FROM customer_revenue;
```

7. BUSINESS INTELLIGENCE AND KPI DEVELOPMENT

7.1 KPI Framework Design

7.1.1 Metric Definitions and Calculations

-- Core business metrics with Presto

WITH metrics AS (

SELECT

DATE_TRUNC('day', date) as day,

COUNT(DISTINCT user_id) as dau,

COUNT(DISTINCT CASE WHEN event_type = 'purchase' THEN user_id END) as purchasers,

SUM(CASE WHEN event_type = 'purchase' THEN revenue ELSE 0 END) as revenue,

COUNT(CASE WHEN event_type = 'purchase' THEN 1 END) as transactions

FROM events

GROUP BY 1

)

SELECT

day,

dau,

purchasers,

revenue,

transactions,

purchasers * 100.0 / NULLIF(dau, 0) as conversion_rate,

revenue / NULLIF(transactions, 0) as aov,

revenue / NULLIF(purchasers, 0) as arpu

FROM metrics;

7.2 Business Metrics and Analytics

7.2.1 Sales and Revenue Analysis

-- Revenue breakdown with Presto

SELECT

```
product_category,  
region,  
COUNT(DISTINCT order_id) as orders,  
SUM(quantity) as units_sold,  
SUM(revenue) as total_revenue,  
AVG(revenue) as avg_order_value,  
APPROX_PERCENTILE(revenue, 0.5) as median_order_value
```

FROM sales

WHERE order_date >= CURRENT_DATE - INTERVAL '30' DAY

GROUP BY product_category, region

ORDER BY total_revenue DESC;

7.2.2 Customer Analytics and Segmentation

-- RFM (Recency, Frequency, Monetary) Analysis in Presto

WITH rfm AS (

SELECT

```
customer_id,  
DATE_DIFF('day', MAX(order_date), CURRENT_DATE) as recency,  
COUNT(DISTINCT order_id) as frequency,  
SUM(revenue) as monetary
```

```

        FROM orders
        WHERE order_date >= CURRENT_DATE - INTERVAL '1' YEAR
        GROUP BY customer_id
    ),
    rfm_scores AS (
        SELECT
            customer_id,
            NTILE(5) OVER (ORDER BY recency) as r_score,
            NTILE(5) OVER (ORDER BY frequency DESC) as f_score,
            NTILE(5) OVER (ORDER BY monetary DESC) as m_score
        FROM rfm
    )
    SELECT
        customer_id,
        CAST(r_score AS VARCHAR) || CAST(f_score AS VARCHAR) || CAST(m_score AS
        VARCHAR) as rfm_segment,
        CASE
            WHEN r_score >= 4 AND f_score >= 4 AND m_score >= 4 THEN 'Champions'
            WHEN r_score >= 3 AND f_score >= 3 AND m_score >= 3 THEN 'Loyal Customers'
            WHEN r_score >= 4 AND f_score <= 2 THEN 'New Customers'
            WHEN r_score <= 2 AND f_score >= 3 THEN 'At Risk'
            ELSE 'Other'
        END as segment
    FROM rfm_scores;

```

7.2.3 Marketing Attribution and Campaign Analysis

```
-- Multi-touch attribution (first-click within 7 days) in Presto

WITH first_touch AS (
    SELECT
        user_id,
        campaign_id,
        MIN(touch_time) as first_touch_time
    FROM marketing_touches
    GROUP BY user_id, campaign_id
),

conversions AS (
    SELECT
        c.user_id,
        c.conversion_time,
        c.revenue,
        ft.campaign_id,
        ft.first_touch_time
    FROM conversions c
    JOIN first_touch ft ON c.user_id = ft.user_id
    WHERE c.conversion_time BETWEEN ft.first_touch_time
        AND ft.first_touch_time + INTERVAL '7' DAY
)

SELECT
    campaign_id,
    COUNT(DISTINCT user_id) as attributed_conversions,
    SUM(revenue) as attributed_revenue,
    AVG(revenue) as avg_conversion_value
```

```
FROM conversions  
GROUP BY campaign_id  
ORDER BY attributed_revenue DESC;
```

8. ADVANCED ANALYTICAL TECHNIQUES

8.1 Statistical Analysis with SQL

8.1.1 Descriptive Statistics and Percentiles

```
-- Comprehensive statistics in Presto
```

```
SELECT  
    COUNT(*) as count,  
    AVG(value) as mean,  
    STDDEV(value) as std_dev,  
    VARIANCE(value) as variance,  
    MIN(value) as min,  
    MAX(value) as max,  
    APPROX_PERCENTILE(value, 0.25) as q1,  
    APPROX_PERCENTILE(value, 0.50) as median,  
    APPROX_PERCENTILE(value, 0.75) as q3,  
    APPROX_PERCENTILE(value, ARRAY[0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]) as deciles  
FROM data_points;
```

8.1.3 Outlier Detection and Anomaly Analysis

```
-- Z-score based outlier detection in Presto
```

```
WITH stats AS (
```

```

SELECT
    AVG(value) as mean,
    STDDEV(value) as std_dev
FROM transactions
),
z_scores AS (
    SELECT
        transaction_id,
        value,
        (value - s.mean) / NULLIF(s.std_dev, 0) as z_score
    FROM transactions t
    CROSS JOIN stats s
)
SELECT * FROM z_scores
WHERE ABS(z_score) > 3; -- Outliers beyond 3 standard deviations

```

8.3 Experiment Analysis and A/B Testing

8.3.2 Control and Treatment Group Analysis

```

-- A/B test analysis with Presto
WITH experiment_results AS (
    SELECT
        variant_group,
        COUNT(DISTINCT user_id) as users,
        COUNT(DISTINCT CASE WHEN converted = 1 THEN user_id END) as converters,
        SUM(revenue) as total_revenue
)
```

```

FROM ab_test_data
WHERE experiment_id = 'exp_123'
GROUP BY variant_group
)
SELECT
variant_group,
users,
converters,
converters * 100.0 / NULLIF(users, 0) as conversion_rate,
total_revenue / NULLIF(users, 0) as revenue_per_user,
total_revenue / NULLIF(converters, 0) as revenue_per_converter
FROM experiment_results
ORDER BY variant_group;

```

9. DATA MODELING FOR ANALYTICS

9.1 Dimensional Modeling Concepts

9.1.1 Star Schema and Snowflake Schema

-- Creating fact and dimension tables in Presto/Hive

```

CREATE TABLE IF NOT EXISTS fact_sales (
    sale_id BIGINT,
    date_key INTEGER,
    product_key INTEGER,
    customer_key INTEGER,
    quantity INTEGER,
    revenue DECIMAL(10,2)

```

```
) WITH (
    format = 'ORC',
    partitioned_by = ARRAY['date_key']
);
```

```
CREATE TABLE IF NOT EXISTS dim_product (
    product_key INTEGER,
    product_id VARCHAR,
    product_name VARCHAR,
    category VARCHAR,
    subcategory VARCHAR
) WITH (
    format = 'ORC'
);
```

10. PERFORMANCE OPTIMIZATION AND SCALABILITY

10.1 Query Performance Tuning

10.1.1 Execution Plan Analysis

```
-- View execution plan in Presto
EXPLAIN (TYPE DISTRIBUTED)
SELECT c.customer_name, SUM(o.total_amount)
FROM customers c
JOIN orders o ON c.customer_id = o.customer_id
GROUP BY c.customer_name;
```

```
-- Analyze query with statistics  
EXPLAIN ANALYZE  
  
SELECT * FROM large_table  
  
WHERE partition_column = '2024-01-01';
```

10.1.2 Query Optimization Techniques

Presto-Specific Best Practices:

1. Use partition predicates to limit data scanned
2. Filter before joining when possible
3. Use appropriate file formats (ORC, Parquet)
4. Leverage bucketing for joins
5. Use APPROX functions for large aggregations

```
-- Optimized query with partition pruning  
  
SELECT *  
  
FROM events  
  
WHERE event_date = DATE '2024-01-01' -- Partition predicate  
  
AND user_id IN (SELECT user_id FROM active_users);
```

-- Use APPROX functions for performance

```
SELECT  
  
APPROX_DISTINCT(user_id) as unique_users,  
  
APPROX_PERCENTILE(revenue, 0.5) as median_revenue  
  
FROM large_table;
```

10.2 Big Data SQL Techniques

10.2.2 Partitioning and Bucketing Strategies

```

-- Create partitioned and bucketed table

CREATE TABLE events_optimized
WITH (
    format = 'ORC',
    partitioned_by = ARRAY['event_date'],
    bucketed_by = ARRAY['user_id'],
    bucket_count = 100
)
AS SELECT * FROM events;

-- Query with partition pruning

SELECT COUNT(*)
FROM events_optimized
WHERE event_date BETWEEN DATE '2024-01-01' AND DATE '2024-01-31';

```

11. MACHINE LEARNING DATASET PREPARATION

11.1 Feature Engineering with SQL

11.1.1 Creating Model-Ready Datasets

```

-- Create features for ML model in Presto

WITH user_features AS (
    SELECT
        user_id,
        COUNT(DISTINCT order_id) as total_orders,
        AVG(order_value) as avg_order_value,
        MAX(order_date) as last_order_date,

```

```

        DATE_DIFF('day', MIN(order_date), MAX(order_date)) as customer_lifespan,
        SUM(CASE WHEN product_category = 'Electronics' THEN 1 ELSE 0 END) as
electronics_orders,
        STDDEV(order_value) as order_value_variance
    FROM orders
    GROUP BY user_id
)
SELECT
    uf.*,
    CASE WHEN total_orders > 10 THEN 1 ELSE 0 END as high_value_customer -- Target
FROM user_features uf;

```

12. ADVANCED SQL PATTERNS FOR ANALYSTS

12.2 Funnel Analysis and User Journey

12.2.1 Multi-step Process Analysis

```

-- E-commerce funnel analysis with Presto

WITH funnel_steps AS (
    SELECT
        user_id,
        MAX(CASE WHEN event_type = 'view_product' THEN 1 ELSE 0 END) as view_product,
        MAX(CASE WHEN event_type = 'add_to_cart' THEN 1 ELSE 0 END) as add_to_cart,
        MAX(CASE WHEN event_type = 'checkout' THEN 1 ELSE 0 END) as checkout,
        MAX(CASE WHEN event_type = 'purchase' THEN 1 ELSE 0 END) as purchase
    FROM user_events
    WHERE event_date >= CURRENT_DATE - INTERVAL '30' DAY
)
```

```

        GROUP BY user_id
    )
SELECT
    COUNT(*) as total_users,
    SUM(view_product) as viewed_product,
    SUM(view_product * add_to_cart) as added_to_cart,
    SUM(view_product * add_to_cart * checkout) as checked_out,
    SUM(view_product * add_to_cart * checkout * purchase) as purchased,
    -- Conversion rates
    SUM(view_product * add_to_cart) * 100.0 / NULLIF(SUM(view_product), 0) as
view_to_cart_rate,
    SUM(view_product * add_to_cart * checkout) * 100.0 /
    NULLIF(SUM(view_product * add_to_cart), 0) as cart_to_checkout_rate,
    SUM(view_product * add_to_cart * checkout * purchase) * 100.0 /
    NULLIF(SUM(view_product * add_to_cart * checkout), 0) as checkout_to_purchase_rate
FROM funnel_steps;

```

12.4 Top-N Queries and Rankings

```

-- Top 3 products per category with tie handling
WITH ranked_products AS (
    SELECT
        category,
        product_name,
        revenue,
        DENSE_RANK() OVER (PARTITION BY category ORDER BY revenue DESC) as rank
    FROM product_sales
)
```

```
)  
SELECT * FROM ranked_products  
WHERE rank <= 3;  
  
-- Using array_agg for top-n collection  
  
SELECT  
    category,  
    SLICE(ARRAY_AGG(product_name ORDER BY revenue DESC), 1, 3) as top_3_products  
FROM product_sales  
GROUP BY category;
```

Key Presto/Trino SQL Patterns and Solutions

Pattern 1: Finding Duplicates

```
-- Method 1: GROUP BY with HAVING  
  
SELECT email, COUNT(*) as duplicate_count  
FROM users  
GROUP BY email  
HAVING COUNT(*) > 1;
```

```
-- Method 2: Window Function
```

```
WITH duplicates AS (  
    SELECT *, ROW_NUMBER() OVER (PARTITION BY email ORDER BY user_id) as rn  
    FROM users  
)  
SELECT * FROM duplicates WHERE rn > 1;
```

Pattern 2: Nth Highest Value

```
-- Method 1: DENSE_RANK

WITH ranked_salaries AS (
    SELECT salary, DENSE_RANK() OVER (ORDER BY salary DESC) as rank
    FROM employees
)

SELECT DISTINCT salary FROM ranked_salaries WHERE rank = 3;
```

-- Method 2: Using OFFSET

```
SELECT DISTINCT salary
FROM employees
ORDER BY salary DESC
LIMIT 1 OFFSET 2;
```

Pattern 3: Consecutive Events

```
-- Find users with 3+ consecutive days of activity using Presto

WITH daily_activity AS (
    SELECT
        user_id,
        activity_date,
        activity_date - INTERVAL '1' DAY * ROW_NUMBER() OVER (
            PARTITION BY user_id ORDER BY activity_date
        ) as group_id
    FROM user_activities
```

```

),
consecutive_groups AS (
    SELECT
        user_id,
        MIN(activity_date) as start_date,
        MAX(activity_date) as end_date,
        COUNT(*) as consecutive_days
    FROM daily_activity
    GROUP BY user_id, group_id
)
SELECT DISTINCT user_id
FROM consecutive_groups
WHERE consecutive_days >= 3;

```

Pattern 4: Sessionization (Time-based grouping)

```

-- 30-minute session windows in Presto
WITH session_markers AS (
    SELECT
        user_id,
        event_time,
        LAG(event_time) OVER (PARTITION BY user_id ORDER BY event_time) as
        prev_event_time,
    CASE
        WHEN DATE_DIFF('minute',
            LAG(event_time) OVER (PARTITION BY user_id ORDER BY event_time),
            event_time) > 30

```

```
        OR LAG(event_time) OVER (PARTITION BY user_id ORDER BY event_time) IS NULL
    THEN 1
    ELSE 0
END as new_session
FROM events
),
session_labels AS (
SELECT
    user_id,
    event_time,
    SUM(new_session) OVER (
        PARTITION BY user_id
        ORDER BY event_time
        ROWS UNBOUNDED PRECEDING
    ) as session_id
FROM session_markers
)
SELECT
    user_id,
    session_id,
    MIN(event_time) as session_start,
    MAX(event_time) as session_end,
    COUNT(*) as event_count,
    DATE_DIFF('minute', MIN(event_time), MAX(event_time)) as session_duration_minutes
FROM session_labels
GROUP BY user_id, session_id;
```

Pattern 5: Working with Arrays and JSON (Presto-specific)

```
-- Array operations

SELECT
    user_id,
    ARRAY_AGG(DISTINCT product_id) as products_purchased,
    CARDINALITY(ARRAY_AGG(DISTINCT product_id)) as unique_products,
    ARRAY_JOIN(ARRAY_AGG(product_name ORDER BY purchase_date), ', ') as product_list
FROM purchases
GROUP BY user_id;

-- JSON operations

SELECT
    JSON_EXTRACT_SCALAR(metadata, '$.utm_source') as utm_source,
    JSON_EXTRACT(metadata, '$.properties') as properties,
    CAST(JSON_EXTRACT(metadata, '$.user_id') AS BIGINT) as user_id
FROM events
WHERE JSON_EXTRACT_SCALAR(metadata, '$.event_type') = 'purchase';
```

Performance Best Practices for Presto/Trino

1. Query Optimization Checklist:

- Use partition predicates to limit data scanned
- Push filters down as early as possible
- Use columnar formats (ORC, Parquet)
- Leverage bucketing for large joins

- Use APPROX functions for large-scale aggregations
- Avoid SELECT * in production queries

2. Join Optimization:

- Put smaller table on the right side of the join
- Use broadcast joins for small dimension tables
- Ensure join keys have matching data types
- Consider bucketing tables on join keys

3. Aggregation Optimization:

- Use APPROX_DISTINCT instead of COUNT(DISTINCT) for large datasets
- Pre-aggregate data where possible
- Use partitioning to reduce data scanned

4. Presto-Specific Tips:

- Use `TRY` functions to handle potential errors gracefully
- Leverage `UNNEST` for array operations
- Use `VALUES` for creating inline tables
- Take advantage of lambda functions for complex array operations

Common Interview Tips for Presto/Trino Environments

1. Key Differences to Remember:

- No stored procedures or triggers
- Limited support for UPDATE/DELETE
- CTEs instead of temp tables
- APPROX functions for better performance
- Strong array and JSON support

2. Performance Considerations:

- Always mention partition pruning
- Discuss file format choices (ORC vs Parquet)
- Consider memory limits for large sorts/aggregations
- Understand distributed query execution

3. Business Context at Scale:

- DAU/MAU calculations with APPROX_DISTINCT
- Sessionization at scale with proper windowing
- Funnel analysis with large event datasets
- Cohort retention with partitioned data

SQL Applied

SQL APPLIED

FIELD CASES

@Airbnb

Realistic Database Schema:

- **Users table:** Tracks both hosts and guests with superhost status
- **Listings:** Includes property types, neighborhoods, instant booking, and business travel readiness
- **Bookings:** Captures guest stays, cancellations, and business travel flags
- **Reviews:** Two-way review system with detailed rating categories (cleanliness, communication, value, etc.)
- **Search logs:** Tracks the user journey from search to booking
- **Calendar:** Dynamic pricing and availability management

Airbnb Business Problems Covered:

1. **Marketplace Analytics:** Supply-demand balance, market saturation, geographic expansion opportunities
2. **Host Performance:** Superhost impact analysis, revenue optimization, pricing strategies
3. **Guest Behavior:** Cohort retention, booking patterns, lifetime value segmentation
4. **Conversion Funnel:** Search to booking conversion, instant booking impact
5. **Trust & Safety:** Review analysis, rating distributions, quality metrics
6. **Revenue Optimization:** Dynamic pricing, seasonal patterns, occupancy rates

Real Airbnb Metrics:

- Booking conversion rates and lead times
- Host lifetime value and performance rankings
- Guest retention by signup cohort

- Search effectiveness and click-through rates
- Pricing optimization based on occupancy
- Superhost vs regular host performance comparison

The questions progress from basic queries about listings and availability to complex problems Airbnb's data team actually faces, such as:

- Identifying high-opportunity markets with unmet demand
 - Analyzing the business impact of instant booking
 - Measuring seasonal pricing effectiveness
 - Building guest loyalty segmentation
 - Tracking conversion through the entire user journey
-

SQL Question Types Source 1

Term Definition Questions

Trigger: a procedure stored within a database, which automatically happens whenever a specific event occurs.

Index: a special lookup table within a database to increase data retrieval speed.

Cursor: a pointer, or identifier, associated with a single or group of rows.

Constraints: rules used to limit the type of data allowed within a table. Common constraints include primary key, foreign key, unique key, and NOT NULL.

ETL (Extract, transform, and load): a data integration process used to combine multiple data sources into one data store, such as a [data warehouse](#).

Primary key, foreign key, and unique key: constraints used to identify records within a table.

Normalization vs. denormalization: techniques used to either divide data into multiple tables to achieve integrity ("normalization") or combine data into a table to increase the speed of data retrieval ("denormalization").

RDBMS vs. DBMS: two types of database management systems. Within a relational database management system (RDBMS) data is stored as a table, while in a database management system (DBMS) its stored as a file.

Clustered vs. non-clustered index: two types of indices used to sort and store data. A clustered index sorts data based on their key values, while a non-clustered index stores data and their records in separate locations.

Comprehension Questions Source 1

What is the purpose of an index in a table? Explain the different types.

What are the types of joins in SQL?

What is the difference between DROP, TRUNCATE, and DELETE statements?

How do you use a cursor?

What is the difference between a HAVING clause and a WHERE clause?

Read more: [SQL vs. MySQL: Differences, Similarities, Uses, and Benefits](#)

Questions about a query

This second category gives you an SQL query and asks you a question about it. This tests your ability to read, interpret, analyze, and debug code written by others.

Forms query analysis questions may take:

Given a query,

Put the clauses in order by how SQL would run them.

Identify the error and correct it.

Predict what the query will return.

Explain what problem the query is meant to solve.

Learn more: [SQL vs. NoSQL: The Differences Explained + When to Use Each](#)

Write a query

Categorization, aggregation, and ratio (CASE, COUNT, or SUM, numerator and denominator)

Joining two tables (JOIN inner vs. left or right)

Modifying a database (INSERT, UPDATE, and DELETE)

Comparison operators (Less than, greater than, equal to)

Organizing data (ORDER BY, GROUP BY, HAVING)

Subqueries

Forms query-writing questions may take:

Given a table or tables with a few sample rows,

List the three stores with the highest number of customer transactions.

Extract employee IDs for all employees who earned a three or higher on their last performance review.

Calculate the average monthly sales by product displayed in descending order.

Find and remove duplicates in the table without creating another table.

Identify the common records between two tables.

SQL Question Types Source 2

Joins: Master inner, left, right, and full joins.

Aggregations: Know GROUP BY , HAVING , and functions like SUM() , COUNT() , etc.

Window Functions: Focus on ROW_NUMBER() , RANK() , LAG() , LEAD() .

Subqueries: Learn how to handle subqueries within SELECT, WHERE, and FROM.

Common Table Expressions (CTEs): Understand how and when to use them.

Indexes and Performance: Learn indexing strategies and how to optimize query performance.

Data Modeling: Understand normalization, denormalization, and keys.

Complex Queries: Be able to write complex queries combining multiple concepts.

Real-world Scenarios: Be prepared to solve business problems with SQL.

Error Handling: Learn how to debug and fix common SQL issues.

SQL Question Types Source 3

1 Joins (inner, left, self joins) and understanding data relationships

2 Aggregations with GROUP BY , COUNT() , SUM() , AVG()

- 3 Filtering using WHERE , HAVING , and subconditions
 - 4 Window functions like RANK() , ROW_NUMBER() , LAG() and LEAD()
 - 5 Subqueries and CTEs, especially for nested logic or transformations
 - 6 Indexes and performance optimization, especially on large datasets
-

SQL Questions Source 1

- 1.) Aggregation (sum vs. count, avg, etc....)
- 2.) How would Select data from table A that is not in table B (they are looking for NOT EXISTS or a LEFT JOIN scenario here)
- 3.) Union vs. Union all
- 4.) Difference in JOINS (usually a real world example is asked here such as "You have a customers table and order table. What JOIN would you use to find all customers that had orders?"
- 5.) Date manipulation (this is tricky, because each of these companies have asked varying levels of complexity. One question was asked "how to get the previous 6 months worth of data", another asked "How would you convert a DATETIME field to just DATE"
- 6.) Inserting data into an already created table
- 7.) Case statements (the questions were always a bit ambiguous here, but I was asked a case statement question in each interview)
- 8.) Subquery or CTE related questions. They cared less about the answer, but more about how these are actually used
- 9.) How to identify duplicates in a table? What about multiple tables?
- 10.) Difference between WHERE and HAVING.
- 11.) Windows Functions (LAG / LEAD here).

BONUS QUESTIONS (this is a good way to stand out as a Data Analyst): How would you improve query performance / what would you do if a query is running slow? How would you improve Data Quality in this scenario?

SQL Questions Source 2

- 1: Counting duplicate rows
- 2: What are Joins and types of join and how do they work

3. Aggregate functions
4. Window functions [100% certain], even though at the job you'll Google the syntax but not for live coding round. Even theoretical questions like which window function to solve a particular problem.

5. More windows function: Difference between Rank and Dense Rank
6. On more advanced levels can be questions can be about subqueries, how SQL works behind the scenes, how we store data.

SQL Questions Source 3

Easy

Generate a report that shows employees who their manager is. This was a SELF-JOIN and I'm so sick of this interview question lmao

Show the latest used product. This was simply using MAX on a datetime field.

Medium:

Find customers with the highest orders between a date span. This involved CTEs, converting a datetime to date, and a JOIN.

Calculate the change over time of products for a date span. This involved some aggregation, a case statement, CTE, and window functions.

Hard:

Find users who were active for 4 consecutive days on our app. Again, this was more CTES, windows functions, and aggregations. Also using HAVING a lot.

Other:

They asked a bit about my experiences with queries running slow and solutions. They asked a bit about indexing and working with "big data." They asked about how I would ensure results are correct with large sets of data. Finally, they asked a bit about data visualization experience via Tableau.

Leetcode premium SQL top 40 (sorted by hard first)

SCHEMA

listings (listing_id, host_id, city, country, room_type, bedrooms, price, created_at, is_active)

hosts (host_id, name, joined_date, superhost, country, response_rate)

guests (guest_id, name, signup_date, country, verified)

bookings (booking_id, listing_id, guest_id, host_id, checkin_date, checkout_date, booking_date, nights, guests_count, total_price, status)

reviews (review_id, booking_id, listing_id, guest_id, host_id, review_date, rating, review_text)

payments (payment_id, booking_id, amount, payment_date, payment_method, status)

amenities (amenity_id, listing_id, amenity_name)

calendar (listing_id, date, available, price)

search_sessions (session_id, guest_id, search_date, filters, results_count)

messages (message_id, sender_id, receiver_id, booking_id, sent_date, message_text)

PRACTICE PROBLEMS

SET 1

QUESTION INDEX OVERVIEW

150 Easy questions (30%): Basic SELECT, filtering, simple joins, basic aggregations

250 Medium questions (50%): Window functions, complex joins, subqueries, date/time analysis, business metrics

100 Difficult questions (20%): Advanced analytics, multi-step analysis, optimization, segmentation, funnel analysis

Basic SELECT & Filtering

Basic SELECT

Concept: SELECT statement **Difficulty:** Easy

Business Purpose: Retrieve basic listing information for inventory review

Question: Retrieve all listing IDs, cities, and prices from the listings table.

```
SELECT DATE(session_start) AS day, COUNT(DISTINCT user_id) AS dau
FROM sessions
WHERE DATE(session_start) = DATE '2025-08-01'
GROUP BY 1;
```

WHERE Clause - Single Condition

Concept: Filtering with WHERE **Difficulty:** Easy

Business Purpose: Find listings in specific market for regional analysis

Question: Find all listings located in San Francisco.

```
SELECT *
FROM listings
WHERE city = 'San Francisco';
```

WHERE Clause - Multiple Conditions

Concept: AND operator **Difficulty:** Easy

Business Purpose: Identify premium active listings for marketing campaigns

Question: Find all active listings in New York with prices greater than \$200 per night.

```
SELECT listing_id, price, room_type  
FROM listings  
WHERE city = 'New York' AND price > 200 AND is_active = 1;
```

...

Concept: ... **Difficulty:** Easy

Business Purpose: ...

Question: ...

```
...
```

Concept: ... **Difficulty:** Easy

Business Purpose: ...

Question: ...

```
...
```

...

Concept: ... **Difficulty:** Easy

Business Purpose: ...

Question: ...

...

Concept: ... **Difficulty:** Easy

Business Purpose: ...

Question: ...

...

Concept: ... **Difficulty:** Easy

Business Purpose: ...

Question: ...

PRACTICE PROBLEMS

SET 2

SQL Applied (Unsorted)

Two-Sided Marketplace Challenge

Scenario: "Host sign-ups have increased 30% quarter-over-quarter, but guest bookings have only grown 10%. What problems could this create, and how would you investigate?"

Structure your approach:

Clarify the Problem: This is a supply-demand imbalance - classic two-sided marketplace challenge.

Potential Issues:

- Lower host occupancy rates → frustrated hosts → churn
- Pricing pressure (more supply = lower prices?)
- Quality dilution if we're accepting lower-quality hosts
- Search experience degradation (too many results?)

Key Metrics to Track:

Host Side:

- Host occupancy rate (% of available nights booked)
- Host earnings per month
- Time to first booking
- Host retention/churn rate
- Host satisfaction scores

Guest Side:

- Search results returned per query
- Booking conversion rate
- Guest satisfaction with listings

Marketplace Health:

- Supply utilization rate
- Take rate (Airbnb's revenue %)
- Market concentration (are bookings concentrated among top hosts?)

Investigative Questions:

Geographic Analysis:

- Is the imbalance uniform across markets?
- Are new hosts in areas with existing high demand or oversupplied markets?

Host Quality:

- What's the quality distribution of new hosts?
- Are they competitive on price/amenities?
- Photo quality, descriptions, response rates?

Guest Demand:

- Has guest search volume increased?
- Are guests finding what they're looking for?
- Has time-to-book changed?

Segment Analysis:

- What types of properties are new hosts adding?
- Is it matching guest search patterns?

Recommendations:

Short-term:

- Pause host acquisition in oversupplied markets
- Improve search ranking for new hosts (cold-start problem)
- Dynamic pricing recommendations for new hosts

Medium-term:

- Focus guest acquisition in undersupplied markets
- Host quality scoring and selective acceptance
- Better host onboarding (photos, pricing, calendars)

Long-term:

- Predictive model: Match host acquisition to demand forecasts
- Market-by-market growth strategies
- Product features to improve utilization (flexible search, last-minute bookings)

Metrics

1. Acquisition Metrics

Metrics focused on attracting new users (guests or hosts) and evaluating channel effectiveness.

- New User Signups: Number of new guests or hosts registering per period (e.g., monthly).
- Customer Acquisition Cost (CAC): Total marketing spend divided by new users acquired.
- Acquisition Channel Efficiency: Performance by channel (e.g., email, social, referrals), measured by cost per acquisition or conversions.
- Referral Rate: Percentage of new signups from referrals.

2. Engagement Metrics

Metrics tracking user interaction with the platform post-acquisition.

- User Engagement Rate: Sessions, searches, or views per active user.
- Messaging Response Rate: Percentage of inquiries responded to by hosts within a timeframe (e.g., 24 hours).
- App/Website Traffic: Unique visitors or page views, segmented by source.
- Active Users: Daily/Monthly Active Users (DAU/MAU) for guests and hosts.

3. Conversion Metrics

Metrics for turning interest into actions, like bookings.

- Conversion Rate: Percentage of searches or inquiries leading to bookings.
- Inquiry-to-Booking Rate: Ratio of host inquiries to completed bookings.
- First Booking Rate: Percentage of new users who book within a set period (e.g., 7 days).

4. Retention Metrics

Metrics assessing user loyalty and repeat behavior.

- Retention Rate: Percentage of users active in subsequent periods (e.g., 30-day or cohort-based).
- Churn Rate: Percentage of users who stop engaging (e.g., no bookings in 90 days).
- Repeat Booking Rate: Percentage of users with multiple bookings.

5. Revenue Metrics

Metrics tied to financial outcomes from marketing efforts.

- Total Revenue: Aggregate booking value or fees collected.
- Average Daily Rate (ADR): Average price per night for booked listings.
- Revenue Per Available Rental (RevPAR): Revenue divided by available listing-nights.
- Average Booking Value: Mean value of bookings, segmented by user type or channel.

6. Cost Metrics

Metrics for efficiency of spending.

- Marketing Spend: Total costs on campaigns, ads, etc.
- Cost Per Acquisition (CPA): Similar to CAC but per channel or campaign.
- Return on Investment (ROI): Revenue generated minus costs, divided by costs.

7. Customer Value Metrics

Long-term value derived from users.

- Lifetime Value (LTV): Projected revenue from a user over their lifetime.
- Customer Lifetime Value (CLV): Similar to LTV, often including host contributions.

8. Funnel Metrics

End-to-end user journey analysis.

- Funnel Drop-off Rates: Percentage lost at each stage (e.g., signup -> search -> booking -> review).
- Booking Lead Time: Average days from search/inquiry to booking.
- Cohort Analysis: Performance by user groups (e.g., signup month) across funnel stages.

9. User Experience Metrics

Metrics reflecting satisfaction and quality.

- Net Promoter Score (NPS): Likelihood of recommendation from surveys.
- Review Scores: Average ratings for listings or hosts.
- Guest/Host Satisfaction Score: Composite from reviews and feedback.
- Superhost Ratio: Percentage of hosts meeting elite criteria (e.g., high ratings, low cancellations).

10. Operational Metrics (Marketing-Adjacent)

Metrics influencing marketing strategies, like supply-demand balance.

- Occupancy Rate: Percentage of available nights booked.
- Average Length of Stay (ALOS): Mean days per booking.
- Year-over-Year Growth: Increases in bookings, revenue, or users.
- Market Penetration: Listings or bookings per city/region.

SQL Applied 1 (Syntax)

SQL Applied 1 (Syntax)

Interview-Length SQL Combinations for Marketing Analytics

[@Claude](#)

1. Month-over-Month Growth by Channel

Combination: DATE_TRUNC + LAG + CASE

```
-- Calculate MoM growth and flag declining channels
SELECT
    DATE_TRUNC('month', booking_date) as month,
    marketing_channel,
    COUNT(*) as bookings,
    LAG(COUNT(*)) OVER (PARTITION BY marketing_channel ORDER BY
    DATE_TRUNC('month', booking_date)) as prev_month,
    CASE
        WHEN LAG(COUNT(*)) OVER (PARTITION BY marketing_channel ORDER BY
        DATE_TRUNC('month', booking_date)) IS NULL THEN NULL
        ELSE (COUNT(*) - LAG(COUNT(*)) OVER (PARTITION BY marketing_channel
        ORDER BY DATE_TRUNC('month', booking_date))) * 100.0 /
            LAG(COUNT(*)) OVER (PARTITION BY marketing_channel ORDER BY
            DATE_TRUNC('month', booking_date))
    END as mom_growth_pct
FROM bookings
WHERE booking_date >= '2024-01-01'
GROUP BY month, marketing_channel
ORDER BY marketing_channel, month;
```

2. Find Returning Users Within Time Window

Combination: SELF JOIN + DATE_DIFF + WHERE

```
-- Users who booked again within 60 days
SELECT
    b1.user_id,
    b1.booking_date as first_booking,
    b2.booking_date as second_booking,
    DATE_DIFF('day', b1.booking_date, b2.booking_date) as days_between,
    b1.marketing_channel as first_channel,
    b2.marketing_channel as second_channel
FROM bookings b1
JOIN bookings b2
    ON b1.user_id = b2.user_id
    AND b2.booking_date > b1.booking_date
    AND DATE_DIFF('day', b1.booking_date, b2.booking_date) <= 60
WHERE b1.booking_date >= '2024-01-01';
```

3. Top Performers by Segment

Combination: CTE + RANK + HAVING

```
-- Top 3 campaigns per channel with minimum spend
WITH campaign_performance AS (
    SELECT
        marketing_channel,
        campaign_id,
        SUM(spend) as total_spend,
        SUM(revenue) as total_revenue,
        SUM(revenue) / NULLIF(SUM(spend), 0) as roas
    FROM marketing_data
    WHERE campaign_date >= '2024-01-01'
    GROUP BY marketing_channel, campaign_id
    HAVING SUM(spend) >= 5000
)
SELECT
    marketing_channel,
    campaign_id,
```

```

total_spend,
total_revenue,
roas,
    RANK() OVER (PARTITION BY marketing_channel ORDER BY roas DESC) as rank
FROM campaign_performance
QUALIFY rank <= 3;

```

4. Cohort Week 1 Retention

Combination: CTE + DATE_TRUNC + DATE_DIFF + CASE

```

-- Week 1 retention rate by signup channel
WITH signups AS (
    SELECT
        user_id,
        DATE_TRUNC('week', signup_date) as cohort_week,
        marketing_channel
    FROM users
    WHERE signup_date >= '2024-01-01'
)
SELECT
    s.cohort_week,
    s.marketing_channel,
    COUNT(DISTINCT s.user_id) as signups,
    COUNT(DISTINCT CASE
        WHEN DATE_DIFF('day', s.cohort_week, b.booking_date) <= 7
        THEN b.user_id
    END) as week1_bookers,
    COUNT(DISTINCT CASE
        WHEN DATE_DIFF('day', s.cohort_week, b.booking_date) <= 7
        THEN b.user_id
    END) * 100.0 / COUNT(DISTINCT s.user_id) as week1_retention_pct
FROM signups s
LEFT JOIN bookings b ON s.user_id = b.user_id
GROUP BY s.cohort_week, s.marketing_channel;

```

5. Revenue Percentiles by Channel

Combination: PERCENTILE_CONT + CASE + WHERE IN

...

6. Attribution Window Analysis

Combination: CTE + DATE_DIFF + BETWEEN + SUM with CASE

```
-- Revenue attributed in different time windows
WITH conversions AS (
    SELECT
        t.marketing_channel,
        t.touch_date,
        b.booking_date,
        b.revenue,
        DATE_DIFF('day', t.touch_date, b.booking_date) as
days_to_conversion
    FROM marketing_touches t
    JOIN bookings b
        ON t.user_id = b.user_id
        AND b.booking_date >= t.touch_date
        AND DATE_DIFF('day', t.touch_date, b.booking_date) <= 30
)
SELECT
    marketing_channel,
    COUNT(*) as total_conversions,
    SUM(CASE WHEN days_to_conversion <= 1 THEN revenue ELSE 0 END) as
day1_revenue,
    SUM(CASE WHEN days_to_conversion <= 7 THEN revenue ELSE 0 END) as
day7_revenue,
    SUM(CASE WHEN days_to_conversion <= 14 THEN revenue ELSE 0 END) as
day14_revenue,
    SUM(revenue) as day30_revenue
FROM conversions
GROUP BY marketing_channel;
```

7. Running Totals with Reset

Combination: SUM OVER + PARTITION BY + DATE_TRUNC

```
-- Cumulative spend per campaign, resetting each quarter
SELECT
    campaign_id,
    DATE_TRUNC('quarter', spend_date) as quarter,
    spend_date,
    daily_spend,
    SUM(daily_spend) OVER (
        PARTITION BY campaign_id, DATE_TRUNC('quarter', spend_date)
        ORDER BY spend_date
    ) as qtd_spend
FROM campaign_spend
WHERE spend_date >= '2024-01-01'
ORDER BY campaign_id, spend_date;
```

8. First and Last Touch Attribution

Combination: CTE + ROW_NUMBER + CASE + WHERE

```
-- Compare first vs last touch channel
WITH ranked_touches AS (
    SELECT
        user_id,
        marketing_channel,
        touch_date,
        ROW_NUMBER() OVER (PARTITION BY user_id ORDER BY touch_date ASC) as
first_touch,
        ROW_NUMBER() OVER (PARTITION BY user_id ORDER BY touch_date DESC)
as last_touch
    FROM marketing_touches
)
SELECT
    first_channel,
    last_channel,
    COUNT(DISTINCT user_id) as users,
    SUM(revenue) as total_revenue
FROM (
    SELECT
```

```

        t1.user_id,
        t1.marketing_channel as first_channel,
        t2.marketing_channel as last_channel,
        b.revenue
    FROM ranked_touches t1
    JOIN ranked_touches t2 ON t1.user_id = t2.user_id
    JOIN bookings b ON t1.user_id = b.user_id
    WHERE t1.first_touch = 1 AND t2.last_touch = 1
) attribution
GROUP BY first_channel, last_channel
ORDER BY total_revenue DESC;

```

9. Active Users by Recency

Combination: MAX + DATE_DIFF + CASE + GROUP BY

```

-- Segment users by last activity
SELECT
    marketing_channel,
    CASE
        WHEN DATE_DIFF('day', MAX(booking_date), CURRENT_DATE) <= 30 THEN
        'active'
        WHEN DATE_DIFF('day', MAX(booking_date), CURRENT_DATE) <= 90 THEN
        'at_risk'
        WHEN DATE_DIFF('day', MAX(booking_date), CURRENT_DATE) <= 180 THEN
        'dormant'
        ELSE 'churned'
    END as user_status,
    COUNT(DISTINCT user_id) as users,
    AVG(DATE_DIFF('day', MAX(booking_date), CURRENT_DATE)) as
    avg_days_since_booking
FROM bookings
WHERE marketing_channel IS NOT NULL
GROUP BY marketing_channel, user_status
ORDER BY marketing_channel,
    CASE user_status
        WHEN 'active' THEN 1
        WHEN 'at_risk' THEN 2
        WHEN 'dormant' THEN 3

```

```
    ELSE 4  
END;
```

10. Conversion Funnel with Drop-off

Combination: CTE + LEFT JOIN + COUNT DISTINCT + COALESCE

```
-- Simple funnel analysis  
WITH funnel AS (  
    SELECT DISTINCT user_id, 'impression' as stage FROM impressions WHERE  
date >= '2024-09-01'  
    UNION  
    SELECT DISTINCT user_id, 'click' as stage FROM clicks WHERE date >=  
'2024-09-01'  
    UNION  
    SELECT DISTINCT user_id, 'signup' as stage FROM signups WHERE date >=  
'2024-09-01'  
    UNION  
    SELECT DISTINCT user_id, 'booking' as stage FROM bookings WHERE date >=  
'2024-09-01'  
)  
SELECT  
    'impression' as stage,  
    COUNT(DISTINCT CASE WHEN stage = 'impression' THEN user_id END) as  
users,  
    COUNT(DISTINCT CASE WHEN stage = 'click' THEN user_id END) * 100.0 /  
        COUNT(DISTINCT CASE WHEN stage = 'impression' THEN user_id END) as  
conversion_to_next  
FROM funnel  
UNION ALL  
SELECT  
    'click' as stage,  
    COUNT(DISTINCT CASE WHEN stage = 'click' THEN user_id END) as users,  
    COUNT(DISTINCT CASE WHEN stage = 'signup' THEN user_id END) * 100.0 /  
        COUNT(DISTINCT CASE WHEN stage = 'click' THEN user_id END) as  
conversion_to_next  
FROM funnel  
UNION ALL  
SELECT
```

```

'signup' as stage,
COUNT(DISTINCT CASE WHEN stage = 'signup' THEN user_id END) as users,
COUNT(DISTINCT CASE WHEN stage = 'booking' THEN user_id END) * 100.0 /
    COUNT(DISTINCT CASE WHEN stage = 'signup' THEN user_id END) as
conversion_to_next
FROM funnel;

```

11. Weekend vs Weekday Performance

Combination: EXTRACT + CASE + DATE_TRUNC + WHERE

```

-- Compare weekend vs weekday bookings
SELECT
    marketing_channel,
    CASE
        WHEN EXTRACT(DOW FROM booking_date) IN (0, 6) THEN 'weekend'
        ELSE 'weekday'
    END as day_type,
    COUNT(*) as bookings,
    AVG(revenue) as avg_revenue,
    SUM(revenue) as total_revenue
FROM bookings
WHERE booking_date >= DATE_TRUNC('month', CURRENT_DATE - INTERVAL '3
months')
GROUP BY marketing_channel, day_type
ORDER BY marketing_channel, total_revenue DESC;

```

12. Users with Multiple Channels

Combination: HAVING + COUNT DISTINCT + STRING_AGG

```

-- Find users exposed to multiple marketing channels
SELECT
    user_id,
    COUNT(DISTINCT marketing_channel) as num_channels,
    STRING_AGG(DISTINCT marketing_channel, ', ' ORDER BY marketing_channel)

```

```
as channels,
    MIN(touch_date) as first_touch,
    MAX(touch_date) as last_touch
FROM marketing_touches
WHERE touch_date >= '2024-01-01'
GROUP BY user_id
HAVING COUNT(DISTINCT marketing_channel) >= 3
ORDER BY num_channels DESC;
```

SQL Applied 2 (Syntax)

SQL Applied 2 (Syntax)

Common SQL Patterns

w/ Problem-Solving

[@Grok](#)

Problems require combining patterns like joins, window functions, CTEs, subqueries, and date handling. Provide your SQL query and explain the business insight it derives.

Note: Dates are based on data up to October 05, 2025.

Aggregation

w/ GROUP BY & HAVING

Description: Summarize by categories and filter aggregates post-grouping; identifies high-performing campaign types for reallocating ad budgets in Q4 2025.

Syntax

```
SELECT
    col1,
    AGG(col2) AS agg_col2
FROM table
GROUP BY col1
HAVING <aggregate_condition>;
```

Airbnb

```
SELECT
    campaign_type,
```

```
    AVG(conversion_rate) AS avg_cr
FROM campaigns
WHERE start_date >= DATE '2025-01-01'
GROUP BY campaign_type
HAVING AVG(conversion_rate) > 0.05;
```

Joins

w/ Aggregations

Description: Combine tables and aggregate results; calculates revenue by user country to prioritize international marketing expansions in 2025.

Syntax

```
SELECT
    t1.col,
    AGG(t2.col) AS agg_val
FROM table1 AS t1
JOIN table2 AS t2
    ON t1.id = t2.id
GROUP BY t1.col;
```

Airbnb

```
SELECT
    u.country,
    SUM(b.booking_value) AS total_revenue
FROM users AS u
JOIN bookings AS b
    ON u.user_id = b.user_id
WHERE b.booking_date >= DATE '2025-01-01'
GROUP BY u.country;
```

Subqueries

Filtering or Derived Metrics

Description: Use nested queries for thresholds or comparisons; finds above-average bookers among 2025 signups for targeted retention.

Syntax

```
SELECT
    col
FROM table
WHERE col >
    (SELECT AGG(col) FROM table);
```

GENERAL SUBQUERY TYPES

Scalar Aggregates

Add metrics like total spend, average rating, or max purchase per user.

```
SELECT
    user_id,
    (SELECT COUNT(*) FROM orders o WHERE o.user_id = u.user_id) AS
order_count
FROM users u;
```

Correlated Lookups

Pull in related data from another table based on the current row.

```
SELECT
    product_id,
    (SELECT category_name FROM categories c WHERE c.category_id =
p.category_id) AS category
FROM products p;
```

SCALAR + CORRELATED SUBQUERY TYPES

w/ SELF JOINS

SCALAR

Table: campaign_performance

campaign_id	channel	spend	impressions	clicks	conversions
1	Search	500	10,000	800	40
2	Social	300	8,000	400	20
3	Display	200	12,000	300	10
4	Email	100	2,000	150	15

Scalar w/ SELECT

Use Case: Attach a global benchmark to each row.

```
SELECT
    campaign_id,
    channel,
    spend,
    (SELECT AVG(spend) FROM campaign_performance) AS avg_spend_all
FROM campaign_performance;
```

Result:

campaign_id	channel	spend	avg_spend_all
1	Search	500	275
2	Social	300	275
3	Display	200	275

	4	Email	100	275
--	---	-------	-----	-----

Every row is enriched with the **same scalar value** (global average spend).

Scalar w/ WHERE

Use Case: Filter rows against a single-value condition.

```
SELECT campaign_id, channel, spend
FROM campaign_performance
WHERE spend > (SELECT AVG(spend) FROM campaign_performance);
```

Result:

	campaign_id	channel	spend
	1	Search	500
	2	Social	300

Only campaigns with spend above the **global average (275)** are returned.

Scalar w/ HAVING

Use Case: Filter groups against a global benchmark.

```
SELECT channel, AVG(conversions) AS avg_conversions
FROM campaign_performance
GROUP BY channel
HAVING AVG(conversions) > (
    SELECT AVG(conversions) FROM campaign_performance
);
```

Result:

	channel	avg_conversions
	Search	40
	Social	20

Only channels with average conversions above the **global average (21.25)** are kept.

Scalar Subqueries Utility

- **SELECT** → Enrich rows with a global constant (benchmarking).
- **WHERE** → Filter individual rows against a global constant.
- **HAVING** → Filter aggregated groups against a global constant.

If you're modeling attribution or running CUPED adjustments, **SELECT** subqueries can help attach pre-period metrics, while **WHERE** subqueries can isolate treatment groups or filter out noisy data.

CORRELATED

Correlated w/ *SELECT*

Query: Attach each campaign's **channel-level average spend**.

```
SELECT
    c1.campaign_id,
    c1.spend,
    (SELECT AVG(c2.spend)
     FROM campaign_performance c2
     WHERE c2.channel = c1.channel) AS avg_spend_channel
  FROM campaign_performance c1;
```

Result:

	campaign_id	channel	spend	avg_spend_channel
	1	Search	500	550
	5	Search	600	550
	2	Social	300	275
	6	Social	250	275
	3	Display	200	200

4	Email	100	100
---	-------	-----	-----

Each row gets the **average spend of its own channel** (row-dependent).

Correlated w/ WHERE

Query: Keep only campaigns whose spend is **above their channel's average spend**.

```
SELECT c1.campaign_id, c1.channel, c1.spend
FROM campaign_performance c1
WHERE c1.spend > (
    SELECT AVG(c2.spend)
    FROM campaign_performance c2
    WHERE c2.channel = c1.channel
);
```

Result:

campaign_id	channel	spend
5	Search	600
2	Social	300

Keeps campaigns above the **average spend of their own channel**.

Correlated w/ HAVING

Use Case: Find channels whose average conversions are above the average conversions of campaigns in the *same spend tier*.

```
SELECT c1.channel, AVG(c1.conversions) AS avg_conversions
FROM campaign_performance c1
GROUP BY c1.channel
HAVING AVG(c1.conversions) > (
    SELECT AVG(c2.conversions)
    FROM campaign_performance c2
    WHERE (CASE WHEN c2.spend >= 300 THEN 'High' ELSE 'Low' END) =
        (CASE WHEN AVG(c1.spend) >= 300 THEN 'High' ELSE 'Low' END)
);
```

Result:

channel	avg_conversions
Search	50
Social	19

SUMMARY

- Search (avg 50) beats the **High-spend tier average** (~40).
- Social (avg 19) beats the **Low-spend tier average** (~14).
- Display and Email don't exceed their tier averages.

Correlated Subqueries Utility

- **SELECT** → enrich each row with a context-specific benchmark (e.g., channel average).
- **WHERE** → filter rows relative to their peer group (e.g., above channel average).
- **HAVING** → Filter groups relative to a peer group benchmark (e.g., channel vs. spend-tier average).

w/ NON-SELF JOINS

SCALAR

They can pull a single value from any table, not just the one in the outer query.

Use Case: Attach the global average spend from the campaign_performance table to each row in a completely different users table.

SELECT

```
u.user_id,  
u.country,  
(SELECT AVG(spend) FROM campaign_performance) AS avg_campaign_spend  
FROM users u;
```

Here, the subquery references a different table (campaign_performance), not a self-join.

CORRELATED

They just need to reference a column from the **outer query**.

Use Case: For each user, check if they've booked more than the average spend of their country (outer query = users, subquery = bookings).

```
SELECT u.user_id, u.country
FROM users u
WHERE u.total_spend > (
    SELECT AVG(b.amount)
    FROM bookings b
    WHERE b.country = u.country
);
```

Here, the subquery references a different table (bookings), but it's still correlated because it depends on u.country from the outer query.

Airbnb

- **Scalar** is great for attaching *global KPIs* (e.g., “global average CTR”).
- **Correlated** is great for *contextual comparisons* (e.g., “this campaign vs. its channel average” or “this user’s spend vs. their own historical average”).

Table: monthly_campaign_performance

campaign_id	channel	month	spend	conversions
1	Search	2025-07-01	400	30
1	Search	2025-08-01	500	40
2	Social	2025-07-01	200	15
2	Social	2025-08-01	300	20
3	Display	2025-07-01	250	12
3	Display	2025-08-01	200	10

Query: Compare each campaign’s spend this month vs. last month

We'll use a correlated subquery that looks up the previous month's spend for the same campaign.

```
SELECT
    curr.campaign_id,
    curr.channel,
    curr.month,
    curr.spend AS current_spend,
    (
        SELECT prev.spend
        FROM monthly_campaign_performance prev
        WHERE prev.campaign_id = curr.campaign_id
        AND prev.month = DATEADD('month', -1, curr.month)
    ) AS previous_spend,
    curr.spend - COALESCE((
        SELECT prev.spend
        FROM monthly_campaign_performance prev
        WHERE prev.campaign_id = curr.campaign_id
        AND prev.month = DATEADD('month', -1, curr.month)
    ), 0) AS spend_delta
FROM monthly_campaign_performance curr
WHERE curr.month = '2025-08-01';
```

Result:

campaign_id	channel	month	current_spend	previous_spend	spend_delta
1	Search	2025-08-01	500	400	+100
2	Social	2025-08-01	300	200	+100
3	Display	2025-08-01	200	250	-50

Sample Use Cases

- **Experimentation:** Compare treatment vs. pre-period metrics (CUPED-style adjustments).
- **Retention/Churn:** Track whether a user's or campaign's activity is increasing or declining month-over-month.
- **Budget Optimization:** Spot campaigns that are scaling up or down in spend and conversions.

Common Table Expressions (CTEs)

Readable Complex Queries

Description: Define temporary result sets for multi-step logic; tracks monthly first-bookings for cohort analysis in 2025.

Syntax

```
WITH cte AS (
    SELECT ...
    FROM ...
)
SELECT ...
FROM cte;
```

Airbnb

```
WITH new_users AS (
    SELECT
        user_id,
        MIN(booking_date) AS first_booking
    FROM bookings
    WHERE booking_date >= DATE '2025-01-01'
    GROUP BY user_id
)
SELECT
```

```
DATE_TRUNC('month', first_booking) AS cohort_month,  
COUNT(user_id) AS new_bookers  
FROM new_users  
GROUP BY cohort_month;
```

Window Functions

Ranking or Running Totals

Description: Calculate across row sets without collapsing rows; computes cumulative conversions to evaluate YTD momentum.

Syntax

```
SELECT  
    col,  
    WINDOW_FUNC(expr) OVER (  
        PARTITION BY part_col  
        ORDER BY order_col  
    ) AS win_val  
FROM table;
```

Airbnb

```
SELECT  
    campaign_id,  
    start_date,  
    conversions,  
    SUM(conversions) OVER (  
        ORDER BY start_date  
        ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW  
    ) AS running_total  
FROM campaigns  
WHERE start_date >= DATE '2025-01-01'  
    AND start_date <= DATE '2025-10-05'  
ORDER BY start_date;
```

Self-Joins

Comparisons Over Time or Hierarchies

Description: Compare rows within the same table, e.g., MoM.

Syntax

```
SELECT
    t1.col,
    t2.col
FROM table AS t1
JOIN table AS t2
    ON t1.id = t2.id
    AND <relationship_condition>;
```

Airbnb

Compare monthly revenue MoM from Jan to Oct 2025.

```
SELECT
    m1.month,
    m1.revenue,
    m2.revenue AS prev_revenue,
    (m1.revenue - m2.revenue) / NULLIF(m2.revenue, 0) AS growth_rate
FROM monthly_metrics AS m1
LEFT JOIN monthly_metrics AS m2
    ON m1.month = date_add('month', 1, m2.month)
WHERE m1.month BETWEEN DATE '2025-01-01' AND DATE '2025-10-01';
```

Comparing Booking Behavior Over Time

Use Case: Track how a guest's booking frequency changes month-over-month; helps identify users with increasing or decreasing engagement—useful for churn modeling or targeting re-engagement campaigns.

```
SELECT
    curr.user_id,
```

```

curr.month AS current_month,
curr.booking_count AS current_bookings,
prev.booking_count AS previous_bookings,
curr.booking_count - COALESCE(prev.booking_count, 0) AS booking_delta
FROM
monthly_bookings curr
LEFT JOIN
monthly_bookings prev
ON curr.user_id = prev.user_id AND curr.month = DATE_ADD('month', 1,
prev.month);

```

Hierarchical Comparison: Campaign vs. Sub-Campaign

Use Case: Compare performance of sub-campaigns to their parent campaign; useful for evaluating budget allocation efficiency across nested campaign structures.

```

SELECT
child.campaign_id,
child.name AS sub_campaign,
parent.name AS parent_campaign,
child.spend,
parent.spend AS parent_spend,
ROUND(child.spend / parent.spend, 2) AS spend_ratio
FROM
campaigns child
JOIN
campaigns parent
ON child.parent_campaign_id = parent.campaign_id;

```

Experimentation: CUPED Pre-Period Matching

Use Case: Match each user's pre-period metric to their post-period performance; enables CUPED adjustment or direct lift analysis for experimentation frameworks.

```

SELECT
post.user_id,
post.conversion_rate AS post_conversion,
pre.conversion_rate AS pre_conversion,
post.conversion_rate - pre.conversion_rate AS lift

```

```
FROM
  experiment_results post
JOIN
  experiment_results pre
ON post.user_id = pre.user_id AND pre.period = 'pre' AND post.period =
'post';
```

Pricing Dynamics: Host Price Changes

Use Case: Detect when a host changes their listing price; useful for dynamic pricing analysis or alerting systems.

```
SELECT
  curr.listing_id,
  curr.date AS current_date,
  curr.price AS current_price,
  prev.price AS previous_price,
  curr.price - prev.price AS price_change
FROM
  listing_prices curr
JOIN
  listing_prices prev
  ON curr.listing_id = prev.listing_id AND curr.date = DATE_ADD('day', 1,
prev.date);
```

Dates

Cohort or Time-Based Analysis

Description: Group by periods and compute period metrics; builds signup cohorts for 2025.

Syntax

```
SELECT
  DATE_TRUNC('month', date_col) AS month,
  AGG(col) AS agg_value
FROM table
GROUP BY month;
```

Airbnb

```
SELECT
    DATE_TRUNC('month', signup_date) AS cohort_month,
    COUNT(user_id) AS signups
FROM users
WHERE signup_date BETWEEN DATE '2025-01-01' AND DATE '2025-10-05'
GROUP BY cohort_month
ORDER BY cohort_month;
```

Top-N Queries

w/ LIMIT and ORDER BY

Description: Retrieve top or bottom results; top 10 listings by reviews for featured promotions.

Syntax

```
SELECT
    col
FROM table
ORDER BY sort_col DESC
LIMIT n;
```

Airbnb

```
SELECT
    listing_id,
    AVG(review_score) AS avg_score
FROM reviews
WHERE review_date >= DATE '2025-01-01'
GROUP BY listing_id
ORDER BY avg_score DESC
LIMIT 10;
```

Conditional Aggregations

w/ CASE

Description: Segment metrics inside aggregates; segments conversions by source for 2025 optimization.

Syntax

```
SELECT
    SUM(CASE WHEN condition THEN measure ELSE 0 END) AS seg_val
FROM table;
```

Airbnb

```
SELECT
    campaign_id,
    SUM(CASE WHEN source = 'email' THEN conversions ELSE 0 END)
        AS email_conversions,
    SUM(CASE WHEN source = 'social' THEN conversions ELSE 0 END)
        AS social_conversions
FROM campaigns
WHERE start_date >= DATE '2025-01-01'
GROUP BY campaign_id;
```

Pivot/Unpivot

Reshaping Data

Description: Pivot via CASE for portability; pivots monthly bookings for cross-city comparison.

Syntax

```
SELECT
    col,
    SUM(CASE WHEN pivot_col = val1 THEN agg_unit ELSE 0 END) AS val1,
    SUM(CASE WHEN pivot_col = val2 THEN agg_unit ELSE 0 END) AS val2
FROM table
```

```
GROUP BY col;
```

Airbnb

```
SELECT
    city,
    SUM(CASE
        WHEN DATE_TRUNC('month', booking_date) = DATE '2025-01-01'
        THEN 1 ELSE 0
    END) AS jan_bookings,
    SUM(CASE
        WHEN DATE_TRUNC('month', booking_date) = DATE '2025-02-01'
        THEN 1 ELSE 0
    END) AS feb_bookings
FROM bookings
GROUP BY city;
```

NULLs & COALESCE

Management

Description: Provide defaults for missing values; defaults unknown sources to 'Organic' for attribution.

Syntax

```
SELECT
    COALESCE(col1, default_val) AS col1_filled
FROM table
WHERE col2 IS NOT NULL;
```

Airbnb

```
SELECT
    user_id,
    COALESCE(referral_source, 'Organic') AS source,
    signup_date
FROM users
```

```
WHERE signup_date >= DATE '2025-01-01';
```

String Manipulation

Cleaning or Extraction

Description: Normalize and extract patterns; normalizes emails and surfaces potential duplicates.

Syntax

```
SELECT
    substr(col, start, len) AS part,
    replace(col, find, replace) AS replaced
FROM table;
```

Airbnb

```
SELECT
    user_id,
    LOWER(TRIM(email)) AS cleaned_email,
    COUNT(*) OVER (
        PARTITION BY LOWER(TRIM(email)))
    ) AS dup_count
FROM users
WHERE email LIKE '%airbnb%'
    AND signup_date <= DATE '2025-10-05'
ORDER BY cleaned_email;
```

SQL Problem Sets for Airbnb Marketing Analytics Interviews

Below is a set of 10 medium-to-hard SQL problems focused on Airbnb's marketing metrics.

These are designed to mimic interview questions, emphasizing analytical thinking, query optimization, and business insights (e.g., user acquisition, retention, campaign ROI, segmentation, and funnel analysis).

Schema

- **users** (user_id, signup_date, country, referral_source, is_host)
- **bookings** (booking_id, user_id, listing_id, booking_date, booking_value, stay_length_days)
- **campaigns** (campaign_id, start_date, end_date, source, conversions, cost)
- **reviews** (review_id, listing_id, user_id, review_date, review_score)
- **listings** (listing_id, host_id, city, price_per_night)
- **monthly_metrics** (month DATE, revenue, bookings_count, active_users)

Problem Set

Problems require combining patterns like joins, window functions, CTEs, subqueries, and date handling. Provide your SQL query and explain the business insight it derives.

Note: Dates are based on data up to October 05, 2025.

Problem 1 (Medium): Cohort Retention Rate

Calculate the 30-day retention rate for users who signed up in each month of 2025. Retention is defined as users who made at least one booking within 30 days of signup. Output: cohort_month, signups, retained_users, retention_rate (as percentage).

Problem 2 (Medium): Campaign ROI by Source

For campaigns running in 2025, compute ROI (revenue generated minus cost, divided by cost) per source. Join with bookings to attribute revenue where booking_date is between campaign start/end and referral_source matches campaign source. Filter for ROI > 1.5. Output: source, total_cost, total_revenue, roi.

Problem 3 (Hard): User Lifetime Value (LTV) Segmentation

Using a CTE, segment users into 'High', 'Medium', 'Low' LTV based on total booking_value (High: > \$5000, Medium: \$1000-5000, Low: < \$1000). Then, calculate average review_score per segment for bookings in 2025. Output: ltv_segment, user_count, avg_review_score.

Problem 4 (Medium): Month-over-Month Growth in Active Hosts

Identify active hosts (those with at least one booking in the month) and compute MoM growth percentage in active hosts from January to September 2025. Use self-join on monthly aggregated data. Output: month, active_hosts, prev_active_hosts, mom_growth_pct.

Problem 5 (Hard): Funnel Conversion Analysis

Build a funnel for user journey in Q3 2025: signups -> first booking (within 7 days) -> repeat booking (within 30 days of first). Use window functions to track per user. Output: total_signups, first_bookers, repeat_bookers, first_booking_rate, repeat_rate.

Problem 6 (Medium): Top Cities by Booking Growth

Rank the top 5 cities by year-over-year booking value growth from 2024 to 2025 (assume prior year data available). Use subqueries for YoY comparison. Output: city, bookings_2024, bookings_2025, growth_pct, rank.

Problem 7 (Hard): Attribution Modeling

Attribute bookings to the last campaign a user interacted with before booking (assume a interactions table with user_id, campaign_id, interaction_date). Use ROW_NUMBER() to get the last interaction per booking. Output: campaign_id, attributed_bookings, total_value.

Problem 8 (Medium): Churn Rate by Referral Source

Calculate monthly churn rate for users signed up in 2025: churned if no activity (booking or review) in the last 90 days as of October 05, 2025. Group by referral_source. Output: referral_source, total_users, churned_users, churn_rate.

Problem 9 (Hard): Pivot Table for Seasonal Metrics

Pivot average stay_length_days by city and quarter in 2025 (Q1-Q3). Use CASE for pivoting. Output: city, q1_avg_stay, q2_avg_stay, q3_avg_stay.

Problem 10 (Medium): Anomaly Detection in Conversions

Flag campaigns in 2025 where daily conversions dropped more than 20% from the 7-day moving average. Use LAG and AVG over window. Output: campaign_id, date, conversions, moving_avg, drop_pct (for flagged days).

SQL Applied 1 (PS)

SQL Applied 1 (PS)

[@ChatGPT](#)

Schema

- **users:** user_id, signup_ts, signup_channel, country.
 - Note: one row per user.
- **sessions:** session_id, user_id, start_ts, landing_url, channel, geo
 - Note: channel is first-touch for session.
- **events:** user_id, session_id, event, ts, listing_id.
 - Note: event in {view, wishlist, start_checkout, book}.
- **bookings:** booking_id, user_id, session_id, ts, payout.
 - Note: ts is booking time.
- **ad_impressions:** impression_id, user_id, campaign_id, channel, ts.
 - Note: ad logs.
- **ad_clicks:** click_id, user_id, campaign_id, channel, ts.
 - Note: click logs.
- **channel_spend:** dt, channel, campaign_id, spend.
 - Note: daily spend.
- **emails:** email_id, user_id, campaign, sent_ts, open_ts, click_ts.
 - Note: lifecycle.
- **experiment_assignments:** user_id, variant.

Problem & Solution Sets

1. Theme: CTR/CVR

Question: Daily CTR and CVR by campaign

Key: date_trunc, LEFT JOIN, BETWEEN, NULLIF, INTERVAL

Pattern: Attribution window join. CTR/CVR by day and campaign

SQL:

```
```sql
SELECT
 date_trunc('day', i.ts) AS d,
 i.campaign_id,
 count(*) AS impressions,
 count(c.click_id) AS clicks,
 count(b.booking_id) AS bookings,
 1.0 * count(c.click_id) / count(*) AS ctr,
 1.0 * count(b.booking_id) /
 NULLIF(count(c.click_id), 0) AS cvr
FROM ad_impressions i
LEFT JOIN ad_clicks c
 ON c.user_id = i.user_id
 AND c.campaign_id = i.campaign_id
 AND c.ts BETWEEN i.ts AND i.ts + INTERVAL '1' day
LEFT JOIN bookings b
 ON b.user_id = i.user_id
 AND b.ts BETWEEN i.ts AND i.ts + INTERVAL '7' day
GROUP BY 1, 2;
```
```

2. Theme: ROAS

Question: ROAS by channel per day.

Keys: date_trunc, LEFT JOIN, GROUP BY, NULLIF.

Pattern: Join daily spend to booked revenue.

SQL:

```

SELECT
  s.dt,
  s.channel,
  sum(s.spend) AS spend,
  sum(b.rev) AS revenue,
  1.0 * sum(b.rev) / NULLIF(sum(s.spend), 0) AS roas
FROM channel_spend s
LEFT JOIN (
  SELECT
    date_trunc('day', b.ts) AS dt,
    sess.channel,
    sum(b.payout) AS rev
  FROM bookings b
  JOIN sessions sess
  ON b.session_id = sess.session_id
  GROUP BY 1, 2
) b
  ON s.dt = b.dt
  AND s.channel = b.channel
GROUP BY 1, 2;

```

3. Theme: Last-touch

Question: Last channel before booking (per booking).

Key : max_by, inequality join.

Pattern: Last-touch attribution.

SQL:

```

SELECT
  b.booking_id,
  max_by(sess.channel, sess.start_ts) AS last_touch
FROM bookings b
JOIN sessions sess
  ON b.user_id = sess.user_id
  AND sess.start_ts <= b.ts

```

GROUP BY 1;

4. Theme: LTV 90d

Question: 90-day LTV by signup cohort month.

Key : WITH, date_trunc, INTERVAL, SUM.

Pattern: Cohort LTV with fixed horizon.

SQL:

```
```sql
```

```
WITH u AS (
 SELECT
 user_id,
 date_trunc('month', signup_ts) AS cohort
 FROM users
),
r AS (
 SELECT
 u.cohort,
 b.user_id,
 sum(b.payout) AS rev
 FROM u
 JOIN bookings b
 ON u.user_id = b.user_id
 AND b.ts <= u.cohort + INTERVAL '90' day
 GROUP BY 1, 2
)
```

```

)
SELECT
 cohort,
 count(DISTINCT user_id) AS users,
 1.0 * sum(rev) / NULLIF(count(DISTINCT user_id), 0) AS ltv_90d
FROM r
GROUP BY 1
ORDER BY 1;
```

```

5. Theme: Funnel

Question: Session-level funnel rates (view→checkout→book).

Key : CASE, MAX, SUM, NULLIF.

Pattern: Binary step flags then ratios.

SQL:

```

```sql
WITH step AS (
 SELECT
 e.session_id,
 max(CASE WHEN event = 'view' THEN 1 END) AS v,
 max(CASE WHEN event = 'start_checkout' THEN 1 END) AS sc,
 max(CASE WHEN event = 'book' THEN 1 END) AS bk
 FROM events e
 GROUP BY 1
)

```

```

SELECT
 count(*) AS sessions,
 1.0 * sum(v) / count(*) AS p_view,
 1.0 * sum(sc) / NULLIF(sum(v), 0) AS p_sc_given_view,
 1.0 * sum(bk) / NULLIF(sum(sc), 0) AS p_book_given_sc
FROM step;
```

```

6. Theme: Rolling

Question: 7-day rolling bookings per channel.

Key : window SUM with ROWS frame.

Pattern: Rolling aggregates.

SQL:

```
```sql
```

```

SELECT
 date_trunc('day', b.ts) AS d,
 s.channel,
 sum(b.payout) AS rev,
 sum(sum(b.payout)) OVER (
 PARTITION BY s.channel
 ORDER BY date_trunc('day', b.ts)
 ROWS BETWEEN 6 PRECEDING AND CURRENT ROW
) AS rev_7d
FROM bookings b
JOIN sessions s

```

```
ON b.session_id = s.session_id
GROUP BY 1, 2;
```

```
```
```

7. Theme: Retention

Question: Month-N retention from signup.

Key : WITH, date_trunc, DISTINCT, date_diff.

Pattern: Retention matrix by cohort and month N.

SQL:

```
```sql  

WITH a AS (
 SELECT
 user_id,
 date_trunc('month', signup_ts) AS cohort
 FROM users
,
m AS (
 SELECT DISTINCT
 e.user_id,
 date_trunc('month', e.ts) AS active_m
 FROM events e
)
SELECT
 cohort,
 date_diff('month', cohort, active_m) AS m_n,
```

```
count(DISTINCT m.user_id) AS active_users
FROM a
JOIN m
ON a.user_id = m.user_id
GROUP BY 1, 2
ORDER BY 1, 2;
```
```

8. Theme: A/B mean

Question: A/B diff in booking rate with stats.

Key : SUM(CASE), proportions, sqrt.

Pattern: Diff-in-proportions with z-score.

SQL:

```
```sql  
WITH base AS (
 SELECT
 t.variant,
 count(*) AS users,
 sum(CASE WHEN b.booking_id IS NOT NULL THEN 1 ELSE 0 END)
 AS booked
 FROM experiment_assignments t
 LEFT JOIN bookings b
 ON b.user_id = t.user_id
 GROUP BY 1
)
```

```

stats AS (
 SELECT
 variant,
 users,
 booked,
 1.0 * booked / users AS cr
 FROM base
)
SELECT
 sA.cr AS cr_A,
 sB.cr AS cr_B,
 (sB.cr - sA.cr) AS diff,
 (sA.cr * (1 - sA.cr) / sA.users +
 sB.cr * (1 - sB.cr) / sB.users) AS var,
 (sB.cr - sA.cr) / sqrt(var) AS z
FROM stats sA
CROSS JOIN stats sB
WHERE sA.variant = 'A'
 AND sB.variant = 'B';
```

```

9. Theme: CUPED

Question: CUPED-style adjusted mean for A/B.

Key : covar_samp, var_samp, window avg.

Pattern: Variance reduction on revenue.

SQL:

```
```sql
WITH d AS (
 SELECT
 t.variant,
 u.user_id,
 COALESCE(pre.pre_rev, 0.0) AS x,
 COALESCE(post.post_rev, 0.0) AS y
 FROM experiment_assignments t
 JOIN users u
 ON t.user_id = u.user_id
 LEFT JOIN (
 SELECT user_id, sum(payout) AS pre_rev
 FROM bookings
 WHERE ts < date '2025-01-01'
 GROUP BY 1
) pre USING (user_id)
 LEFT JOIN (
 SELECT user_id, sum(payout) AS post_rev
 FROM bookings
 WHERE ts >= date '2025-01-01'
 GROUP BY 1
) post USING (user_id)
),
theta AS (
 SELECT covar_samp(y, x) / NULLIF(var_samp(x), 0) AS th
```

```

FROM d
)
SELECT
variant,
avg(y - th * (x - avg(x) OVER ())) AS cuped_mean
FROM d
CROSS JOIN theta
GROUP BY 1;
```

```

10. Theme: DiD

Question: Geo DiD on bookings before and after launch.

Key : WITH, CASE, GROUP BY.

Pattern: 2×2 cells, compute DiD from aggregates.

SQL:

```

```sql
WITH g AS (
SELECT
geo,
CASE WHEN geo IN ('US_TX', 'US_FL') THEN 1 ELSE 0 END AS trt
FROM sessions
GROUP BY 1, 2
),
ba AS (
SELECT

```

```

CASE WHEN b.ts < date '2025-06-01' THEN 0 ELSE 1 END AS post,
s.geo,
count(*) AS books
FROM bookings b
JOIN sessions s
ON b.session_id = s.session_id
GROUP BY 1, 2
)
SELECT
post, trt, sum(books) AS n
FROM ba
JOIN g USING (geo)
GROUP BY 1, 2;
```

```

11. Theme: Sessionize

Question: Build 30-min sessions from events.

Key : LAG, date_diff, running SUM.

Pattern: Gaps-and-islands.

SQL:

```

```sql
WITH e AS (
SELECT
user_id,
ts,

```

```

CASE
WHEN lag(ts) OVER (
 PARTITION BY user_id ORDER BY ts
) IS NULL
OR date_diff(
 'minute',
 lag(ts) OVER (PARTITION BY user_id ORDER BY ts),
 ts
) > 30
THEN 1 ELSE 0
END AS new_s
FROM events
),
s AS (
SELECT
 user_id,
 ts,
 sum(new_s) OVER (
 PARTITION BY user_id ORDER BY ts
) AS sid
FROM e
)
SELECT * FROM s;
```

```

12. Theme: MTA simple

Question: Equal split across channels in 7-day lookback.

Key : array_agg, array_distinct, UNNEST, cardinality.

Pattern: Simple equal-weight attribution.

SQL:

```
```sql
WITH hist AS (
 SELECT
 b.booking_id,
 b.payout,
 array_distinct(array_agg(s.channel)) AS chs
 FROM bookings b
 JOIN sessions s
 ON b.user_id = s.user_id
 AND s.start_ts BETWEEN b.ts - INTERVAL '7' day AND b.ts
 GROUP BY 1, 2
),
explode AS (
 SELECT
 booking_id,
 payout,
 ch,
 cardinality(chs) AS k
 FROM hist
 CROSS JOIN UNNEST(chs) AS t(ch)
)
```

```
SELECT
 ch AS channel,
 sum(payout * 1.0 / k) AS attrib_rev
FROM explode
GROUP BY 1;
````
```

13. Theme: Percentiles

Question: Time-to-book percentiles.

Key : min_by, date_diff, approx_percentile.

Pattern: Distribution summary.

SQL:

```
```sql
WITH t AS (
 SELECT
 b.user_id,
 min_by(s.start_ts, s.start_ts) AS first_s,
 min_by(b.ts, b.ts) AS first_b
 FROM bookings b
 JOIN sessions s
 ON b.session_id = s.session_id
 GROUP BY 1
)
SELECT
 approx_percentile(
```

```
date_diff('hour', first_s, first_b),
ARRAY[0.5, 0.9, 0.95]
) AS p50_p90_p95
FROM t;
```
```

14. Theme: UTM parse

Question: Extract UTM params from landing_url.

Key : url_extract_parameter.

Pattern: URL parsing for attribution.

SQL:

```
```sql  
SELECT
session_id,
url_extract_parameter(landing_url, 'utm_source') AS utm_source,
url_extract_parameter(landing_url, 'utm_campaign') AS utm_campaign
FROM sessions;
```
```

15. Theme: De-dupe

Question: Pick one device per user by latest seen.

Key : max_by, GROUP BY.

Pattern: Latest per key.

SQL:

```
```sql
SELECT
 user_id,
 max_by(device, start_ts) AS primary_device
FROM sessions
GROUP BY 1;
```

```

16. Theme: Pivot

Question: Pivot revenue by channel (wide).

Key : conditional SUM.

Pattern: Pivot-like conditional aggregation.

SQL:

```
```sql
SELECT
 date_trunc('day', b.ts) AS d,
 sum(CASE WHEN s.channel = 'Paid Search' THEN b.payout END)
 AS paid_search,
 sum(CASE WHEN s.channel = 'Paid Social' THEN b.payout END)
 AS paid_social,
 sum(CASE WHEN s.channel = 'Affiliates' THEN b.payout END)
 AS affiliates
FROM bookings b
JOIN sessions s
 ON b.session_id = s.session_id

```

```
GROUP BY 1
```

```
ORDER BY 1;
```

```
```
```

17. Theme: Top paths

Question: Top 10 channel paths in 7 days pre-booking.

Key : array_agg ORDER BY, array_join.

Pattern: Path sequencing.

SQL:

```
```sql
```

```
WITH seq AS (
```

```
 SELECT
```

```
 b.booking_id,
```

```
 array_join(
```

```
 array_agg(s.channel ORDER BY s.start_ts),
```

```
 ' > '
```

```
) AS path
```

```
 FROM bookings b
```

```
 JOIN sessions s
```

```
 ON b.user_id = s.user_id
```

```
 AND s.start_ts BETWEEN b.ts - INTERVAL '7' day AND b.ts
```

```
 GROUP BY 1
```

```
)
```

```
SELECT
```

```
 path,
```

```
count(*) AS n
FROM seq
GROUP BY 1
ORDER BY n DESC
LIMIT 10;
```
```

18. Theme: Anomaly

Question: Z-score for daily spend by channel.

Key : window AVG, STDDEV_SAMP, ROWS frame.

Pattern: Rolling z-score.

SQL:

```
```sql  
WITH x AS (
 SELECT
 dt,
 channel,
 spend,
 avg(spend) OVER (
 PARTITION BY channel
 ORDER BY dt
 ROWS BETWEEN 14 PRECEDING AND CURRENT ROW
) AS mu,
 stddev_samp(spend) OVER (
 PARTITION BY channel
```

```

 ORDER BY dt
 ROWS BETWEEN 14 PRECEDING AND CURRENT ROW
) AS sd
 FROM channel_spend
)
SELECT
 dt, channel, spend,
 (spend - mu) / NULLIF(sd, 0) AS z
FROM x;
```

```

19. Theme: Dedup clicks

Question: First click after impression within 24h.

Key : LEFT JOIN, BETWEEN, min_by.

Pattern: First-event within window.

SQL:

```

```sql
SELECT
 i.impression_id,
 min_by(c.click_id, c.ts) AS first_click
FROM ad_impressions i
LEFT JOIN ad_clicks c
 ON c.user_id = i.user_id
 AND c.campaign_id = i.campaign_id
 AND c.ts BETWEEN i.ts AND i.ts + INTERVAL '24' hour

```

GROUP BY 1;

```

20. Theme: CPA

Question: CPA by campaign using first conversion in 7 days.

Key : min_by, CASE, NULLIF.

Pattern: Horizon-trimmed first conversion.

SQL:

```sql

WITH conv AS (

SELECT

c.campaign\_id,

c.click\_id,

min\_by(b.booking\_id, b.ts) AS first\_book

FROM ad\_clicks c

LEFT JOIN bookings b

ON b.user\_id = c.user\_id

AND b.ts BETWEEN c.ts AND c.ts + INTERVAL '7' day

GROUP BY 1, 2

),

agg AS (

SELECT

campaign\_id,

sum(CASE WHEN first\_book IS NOT NULL THEN 1 ELSE 0 END)

AS convs

```
FROM conv
GROUP BY 1
)
SELECT
s.campaign_id,
sum(s.spend) / NULLIF(a.convs, 0) AS cpa
FROM channel_spend s
JOIN agg a
ON s.campaign_id = a.campaign_id
GROUP BY 1;
```

```

21. Theme: HLL

Question: Approx unique users per channel.

Key : approx_distinct.

Pattern: Scale-aware uniques.

SQL:

```
```sql
SELECT
s.channel,
approx_distinct(b.user_id) AS users
FROM bookings b
JOIN sessions s
ON b.session_id = s.session_id
GROUP BY 1;
```

```

22. Theme: Cohort ROAS

Question: Signup-month cohort ROAS in 30 days.

Key : WITH, date_trunc, INTERVAL, NULLIF.

Pattern: Cohort revenue vs spend month.

SQL:

```sql

```
WITH c AS (
 SELECT
 user_id,
 date_trunc('month', signup_ts) AS cohort
 FROM users
),
r AS (
 SELECT
 c.cohort,
 sum(b.payout) AS rev
 FROM c
 JOIN bookings b
 ON b.user_id = c.user_id
 AND b.ts <= c.cohort + INTERVAL '30' day
 GROUP BY 1
),
spend AS (
```

```

SELECT
 date_trunc('month', dt) AS m,
 sum(spend) AS spend
FROM channel_spend
GROUP BY 1

)
SELECT
 r.cohort,
 r.rev,
 s.spend,
 1.0 * r.rev / NULLIF(s.spend, 0) AS roas
FROM r
JOIN spend s
 ON r.cohort = s.m;
```

```

23. Theme: Country mix

Question: Share of revenue by country per week.

Key : window SUM, ratio of sums.

Pattern: Mix share by period.

SQL:

```

```sql
SELECT
 date_trunc('week', b.ts) AS wk,
 u.country,

```

```

sum(b.payout) AS rev,
1.0 * sum(b.payout)
/ sum(sum(b.payout)) OVER (
 PARTITION BY date_trunc('week', b.ts)
) AS share

FROM bookings b
JOIN users u
ON b.user_id = u.user_id
GROUP BY 1, 2;
```

```

24. Theme: Email lift

Question: Open→book conversion within 3 days of send.

Key : MAX(CASE), AVG, BETWEEN.

Pattern: Post-exposure outcome window.

SQL:

```

```sql
WITH got_email AS (
 SELECT user_id, sent_ts FROM emails
),
outcome AS (
 SELECT
 g.user_id,
 max(
 CASE

```

```
WHEN b.ts BETWEEN g.sent_ts
AND g.sent_ts + INTERVAL '3' day
THEN 1
END
) AS booked
FROM got_email g
LEFT JOIN bookings b
ON b.user_id = g.user_id
GROUP BY 1
)
SELECT avg(booked) AS conv_rate
FROM outcome;
```

\*\*\*

# **SQL Applied 2 (PS)**

# SQL Applied 2 (PS)

[@Perplexity](#)

## Common Analytics SQL Interview Questions & Solution Patterns

### 1. Aggregation

Typical Questions:

- "What's the total number of bookings per city per month?"
- "Find the average revenue per booking by campaign."

Pattern:

```
SELECT city, MONTH(booking_date) AS month, COUNT(*) AS bookings
FROM bookings
GROUP BY city, MONTH(booking_date)
```

Use GROUP BY for breakdowns, with aggregate functions like COUNT, SUM, AVG.

### 2. Filtering and Sorting

Typical Questions:

- "List bookings for Paris in January, sorted by revenue."

- "Show the top 10 campaigns with the highest conversion rate."

**Pattern:**

```
SELECT *
FROM bookings
WHERE city = 'Paris' AND booking_date BETWEEN '2025-01-01' AND '2025-01-31'
ORDER BY revenue DESC
LIMIT 10
```

Use WHERE for filters, ORDER BY and LIMIT for rankings.

## 3. Joins and Attribution

**Typical Questions:**

- "Total bookings per marketing campaign."
- "Attribute bookings to the last campaign click."

**Pattern:**

```
SELECT mc.campaign_name, COUNT(*) AS total_bookings
FROM bookings b
JOIN marketing_campaigns mc ON b.campaign_id = mc.campaign_id
GROUP BY mc.campaign_name
```

## 4. Window Functions (Rolling/Funnel/Ranking)

**Typical Questions:**

- "7-day rolling average of bookings per city."
- "Rank campaigns by conversions."

**Pattern:**

```
SELECT city, booking_date, bookings,
 AVG(bookings) OVER (PARTITION BY city ORDER BY booking_date ROWS BETWEEN 6
PRECEDING AND CURRENT ROW) AS avg_7d
FROM daily_bookings
```

Use OVER (...) for running, ranking, or partitioned calculations.

## 5. Funnel Analysis/Retention Cohorts

**Typical Questions:**

- "What percent of users made bookings after their first exposure to a campaign?"
- "What is user retention after 30/60/90 days?"

**Pattern:**

- Create first-touch and subsequent event subqueries using MIN(...) or LAG(...) functions, join or filter on intervals.
- Calculate percentages by dividing cohort outcomes by cohort size.

## 6. Case Statements for Segmentation

**Typical Questions:**

- "Classify bookings as 'High', 'Medium', or 'Low' revenue."

**Pattern:**

```
SELECT booking_id,
CASE
 WHEN revenue >= 300 THEN 'High'
 WHEN revenue >= 150 THEN 'Medium'
 ELSE 'Low'
```

```
END AS revenue_segment
```

```
FROM bookings
```

## 7. Missing Data & Time Series Completion

**Typical Questions:**

- "How to calculate moving averages if some dates are missing?"

**Pattern:**

- Generate a calendar table (sequence() in Presto), left join it to your data, fill NULLs with defaults.
- This ensures accurate window functions by date, not just by row count.

## 8. Self-Joins or CTEs for Sequential Analysis

**Typical Questions:**

- "Calculate days between campaign click and booking for each user."
- "Find month-over-month changes."

**Pattern:**

- Use WITH clauses (CTEs) to create intermediate result sets.
- Use LAG() or date arithmetic for calculating differences.

## Schema

### Bookings

booking_id	user_id	city	booking_date	campaign_id	revenue	device_type	rating
------------	---------	------	--------------	-------------	---------	-------------	--------

1	101	Paris	2025-01-15	201	250	mobile	5
2	102	London	2025-01-16	202	300	desktop	4
3	103	New York	2025-02-01	NULL	400	mobile	4
4	104	Paris	2025-02-17	203	225	tablet	3
5	101	Paris	2025-03-05	202	350	mobile	5

## Marketing Campaigns

campaign_id	campaign_name	channel	launch_date	promo_code
201	Winter Blast	email	2025-01-10	WB2025
202	London Calling	social	2025-01-12	LC2025
203	City Escape	search	2025-02-10	CE2025
204	Global Spring	referral	2025-03-01	GS2025

## Users

user_id	user_email	signup_date
101	alice@sample.com	2025-01-05
102	bob@sample.com	2025-01-06
103	carol@sample.com	2025-01-20
104	dave@sample.com	2025-02-14

## Listings

listing_id	host_id	city	active	avg_rating
L10	H201	Paris	TRUE	4.7
L11	H202	London	TRUE	4.8
L12	H203	New York	FALSE	4.3
L13	H204	Paris	TRUE	4.6

## Campaign Clicks

click_id	campaign_id	user_id	click_time	device_type
C101	201	101	2025-01-11	mobile
C102	202	102	2025-01-13	desktop
C103	203	104	2025-02-15	tablet
C104	201	103	2025-01-12	mobile

# Problem Sets

## Easy

### Foundations

1. List all Airbnb bookings in January 2025, sorted by booking date.

```
SELECT booking_id
FROM bookings b
WHERE booking_date BETWEEN '2025-01-01' AND '2025-01-31'
ORDER BY booking_date DESC
```

2. Count the total number of listings in each city.
3. Find the average star rating for each listing.
4. Which users signed up in the past 30 days?
5. How many marketing campaigns ran last quarter?
6. Find the unique number of guests per booking.
7. Count users who have never made a booking.
8. What is the total revenue generated in February 2025?
9. List distinct marketing channels (e.g., email, social, referral) used in the last year.
10. What percentage of bookings were attributed to paid campaigns?

## Medium

### Aggregations, Joins & More

11. For each campaign, show total bookings and average booking value.
12. List the top 5 listings by bookings, including average star rating.
13. Retrieve the number of active listings per host, filtered by a minimum average rating of 4.5.
14. Find the daily number of bookings for each city, showing the 7-day moving average (fill in missing dates!).
15. Calculate click-through rate (CTR) by campaign: ratio of users clicking to users exposed.

16. Show the percentage of first-time bookings coming from each marketing channel.
17. Find users who made bookings from more than one marketing channel.
18. Identify the most common device type used to make bookings.
19. Find the ratio of bookings made using promotions vs. regular price.
20. Which cities have the highest average booking duration?

## Hard

Window Functions, Cohorts, Attribution, Advanced Analytics

21. For each listing, find the month-over-month growth rate in bookings this year.
22. List the top 3 campaigns by conversions and compute their conversion funnel drop-off at each stage (impressions → clicks → bookings).
23. Attribute bookings to the marketing campaign that had the last touch before the booking event, for each user.
24. Compute retention rates: What percentage of users booked more than once within 90 days of their first booking?
25. For each campaign, calculate the median time between first exposure and first booking.
26. Identify properties that saw an increase in bookings after a major event or campaign launch—use window functions.
27. Which listings saw a decline in ratings after a price change promoted through a campaign?
28. Segment Airbnb users into cohorts by signup month and calculate booking rates per cohort over the following 6 months.
29. Analyze the effectiveness of a recent A/B test: Compare booking conversions between test and control groups.
30. Given a table of bookings, campaigns, and user actions, determine which sequence of marketing activities most strongly predicts a booking, using CTEs and aggregation.

# Solution Sets

1. List all Airbnb bookings in January 2025, sorted by booking date.

```
SELECT *
FROM bookings
WHERE booking_date BETWEEN DATE '2025-01-01' AND DATE '2025-01-31'
ORDER BY booking_date;
```

2. Count the total number of listings in each city.

```
SELECT city, COUNT(*) AS num_listings
FROM listings
GROUP BY city;
```

3. Find the average star rating for each listing.

```
SELECT listing_id, AVG(rating) AS avg_rating
FROM bookings
GROUP BY listing_id;
```

(Assumes the bookings table references each listing\_id.)

4. Which users signed up in the past 30 days?

```
SELECT user_id, user_email
FROM users
WHERE signup_date >= CURRENT_DATE - INTERVAL '30' DAY;
```

5. How many marketing campaigns ran last quarter?

```
SELECT COUNT(*) AS num_campaigns
FROM marketing_campaigns
WHERE launch_date BETWEEN DATE '2025-04-01' AND DATE '2025-06-30';
(Adjust dates according to the "last quarter" at your interview date.)
```

**6. Find the unique number of guests per booking.**

```
SELECT booking_id, COUNT(DISTINCT user_id) AS unique_guests
FROM bookings
GROUP BY booking_id;
```

**7. Count users who have never made a booking.**

```
SELECT COUNT(*) AS users_no_booking
FROM users u
LEFT JOIN bookings b ON u.user_id = b.user_id
WHERE b.booking_id IS NULL;
```

**8. What is the total revenue generated in February 2025?**

```
SELECT SUM(revenue) AS feb_revenue
FROM bookings
WHERE booking_date BETWEEN DATE '2025-02-01' AND DATE '2025-02-28';
```

**9. List distinct marketing channels used in the last year.**

```
SELECT DISTINCT channel
FROM marketing_campaigns
WHERE launch_date >= CURRENT_DATE - INTERVAL '1' YEAR;
```

## **10. What percentage of bookings were attributed to paid campaigns?**

```
SELECT
```

```
 100.0 * COUNT(DISTINCT booking_id) FILTER (WHERE campaign_id IS NOT NULL) /
 COUNT(*) AS percent_paid
```

```
FROM bookings;
```

## **11. For each campaign, show total bookings and average booking value.**

```
SELECT c.campaign_id, c.campaign_name, COUNT(b.booking_id) AS total_bookings,
 AVG(b.revenue) AS avg_booking_value
```

```
FROM marketing_campaigns c
```

```
LEFT JOIN bookings b ON c.campaign_id = b.campaign_id
```

```
GROUP BY c.campaign_id, c.campaign_name;
```

## **12. List the top 5 listings by bookings, including average star rating.**

```
SELECT listing_id, COUNT(*) AS total_bookings, AVG(rating) AS avg_rating
```

```
FROM bookings
```

```
GROUP BY listing_id
```

```
ORDER BY total_bookings DESC
```

```
LIMIT 5;
```

## **13. Retrieve the number of active listings per host, filtered by a minimum average rating of 4.5.**

```
SELECT host_id, COUNT(listing_id) AS active_listings
FROM listings
WHERE active = TRUE AND avg_rating >= 4.5
GROUP BY host_id;
```

## **14. Find the daily number of bookings for each city, showing the 7-day moving average (fill in missing dates).**

(Requires generating a date calendar; see previous instructions.)

## **15. Calculate click-through rate (CTR) by campaign.**

sql

```
SELECT
 c.campaign_id,
 COUNT(DISTINCT clk.user_id) AS unique_clicks,
 -- Suppose we have an 'impressions' table or can estimate
 COUNT(DISTINCT clk.user_id) * 1.0 / NULLIF(SUM(impressions), 0) AS ctr
FROM marketing_campaigns c
```

```
LEFT JOIN clicks clk ON c.campaign_id = clk.campaign_id
GROUP BY c.campaign_id;
```

## **16. Show the percentage of first-time bookings coming from each marketing channel.**

```
WITH ft_bookings AS (
 SELECT user_id, MIN(booking_date) AS first_booking
 FROM bookings
 GROUP BY user_id
)

SELECT
 mc.channel,
 COUNT(*) AS num_first_time,
 100.0 * COUNT(*) / (SELECT COUNT(*) FROM ft_bookings) AS pct_first_time

FROM bookings b
JOIN ft_bookings ft ON b.user_id = ft.user_id AND b.booking_date = ft.first_booking
JOIN marketing_campaigns mc ON b.campaign_id = mc.campaign_id
GROUP BY mc.channel;
```